

# Scalable Machine Learning in Python with Dask

10/06/2021

DoD High Performance Computing and Modernization Program (HPCMP)  
User Productivity Enhancement and Training (PET)

# Light Review of terms covered during intro course


- **Dask.array**
  - cuts up large arrays into many small ones.
- **Dask.dataframe**
  - Cuts up large dataframes into many pieces.
- **Dask.Delay**
  - Puts off bringing data into memory, until after task graph is developed.
- **Dask.bag**
  - Used to parallelize computation on unstructured or semi-structured data, like text data.
- **Trask Graphs**




Intro Course on [Dask](#) available

# The Dask Benefits

## Scale Numpy Workflows



## Scale Pandas workflows




- **Parallel computing / multiprocessing :**
  - Use all cores on system (even on laptops).
- **Larger-than-memory problems:**
  - Not limited by large datasets, if using Dask.
  - work on datasets that are larger than your available memory can handle.
  - Chunk/break up your array into many small pieces, and compute on the chunks in parallel.
  - Only stream chunk data from disk when needed for computation.
- **Blocked Algorithms:**
  - Fracture large computations to many smaller computations, and execute on the smaller computations.

# Parallelize Scikit-learn / Joblib

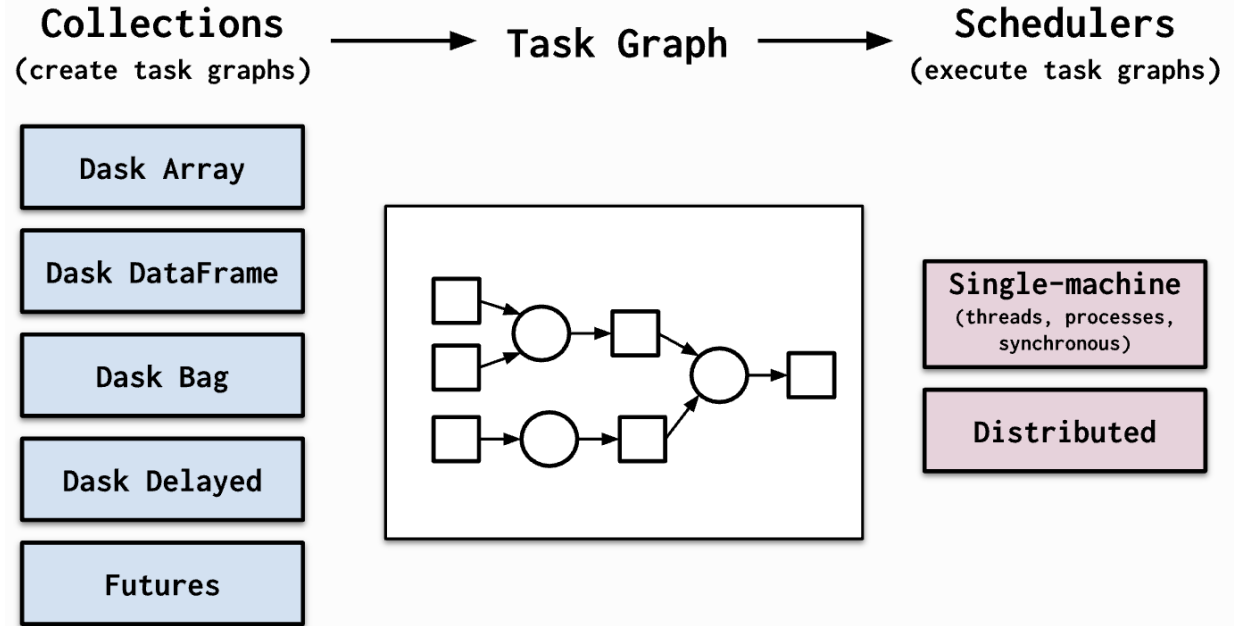
- **Scale Machine Learning APIs**
  - Scikit-learn
  - XGBoost
- **Enable scalable training and predictions on large models & large datasets.**
- **With use of wrappers for Tensorflow & Pytorch Models, integrate Dask with Neural Network models.**
  - We will be proving the concepts with light models.

## scikit-learn

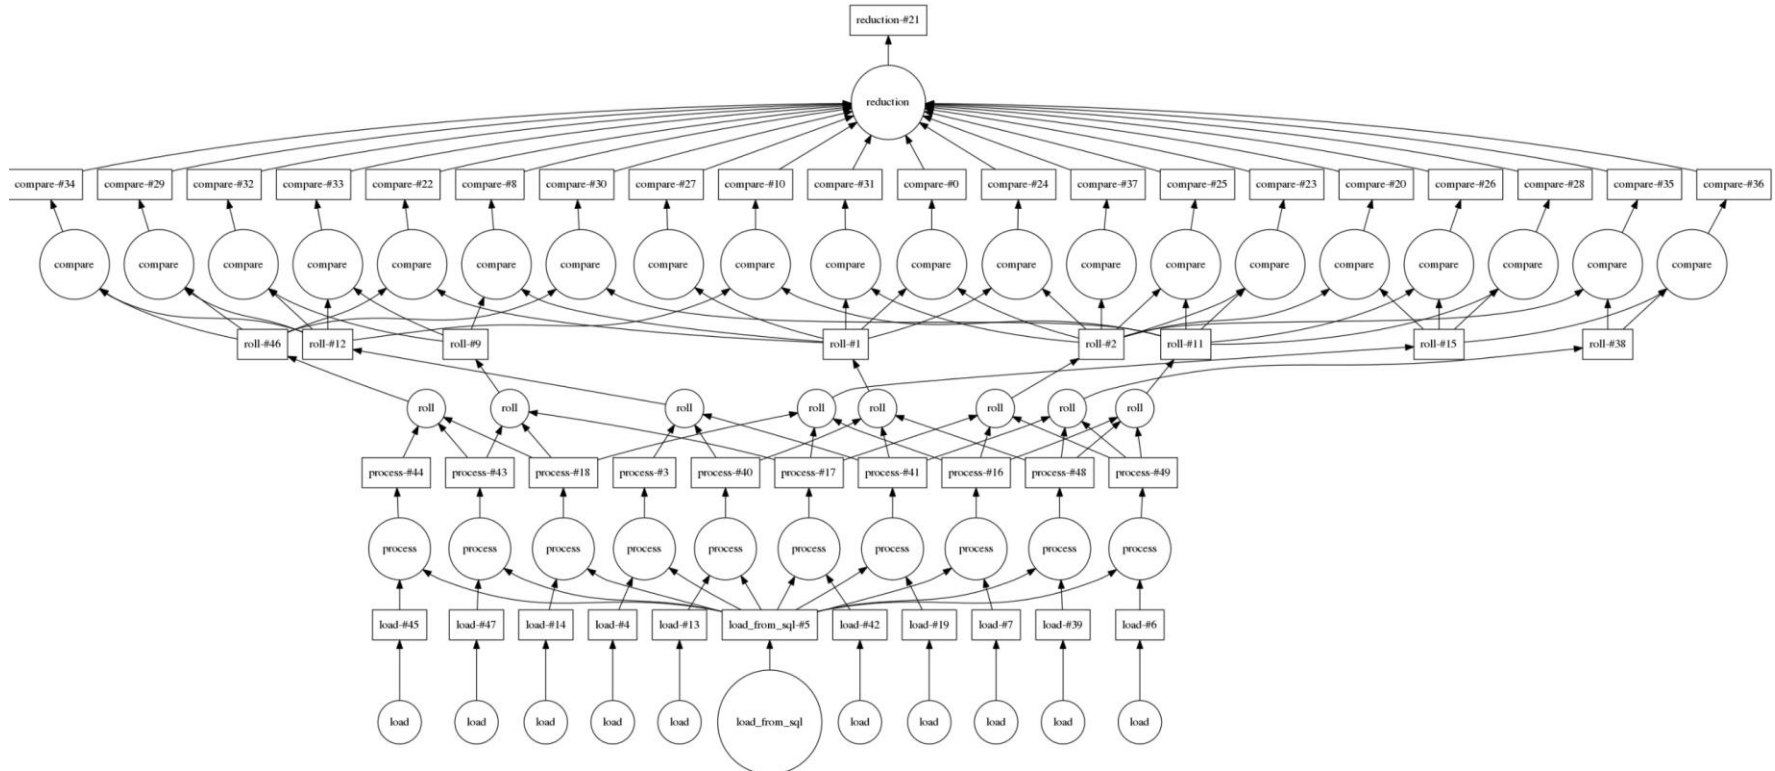


# Use of Task Graphs


- First Generates Task Graph.
- Execute them on parallel nodes.



# Dask Task Graph - Example



# Parallel Computing on Chunks



# Familiar to Python Users

- Designed for Python Ecosystem
- Familiar APIs for Python Users
- Scales upto 1000 node clusters

```
# Arrays implement the NumPy API
import dask.array as da
x = da.random.random(size=(10000, 10000),
                        chunks=(1000, 1000))
x + x.T - x.mean(axis=0)
```

```
# Dataframes implement the pandas API
import dask.dataframe as dd
df = dd.read_csv('s3://.../2018-*-.csv')
df.groupby(df.account_id).balance.sum()
```

```
# Dask-ML implements the scikit-learn API
from dask_ml.linear_model \
    import LogisticRegression
lr = LogisticRegression()
lr.fit(train, test)
```



# Easily Parallelize existing codebases

- **Example of using existing codes:**

- Easily apply lazy function `dask.delayed()` to existing functions to improve computational speeds.
- Computation would not occur until the `dask.compute()` function is executed.

- **Dask is flexible**

- Supports current scikit-learn workflows
- You may also develop your own new scalable algorithms.

```
grid_search.fit(data.data, data.target)
```




```
with joblib.parallel_backend('dask'):  
    grid_search.fit(data.data, data.target)
```

```
func1 = dask.delayed(func1)  
func2 = dask.delayed(func2)  
  
for i in X_list:  
    for j in Y_list:  
        if i > j:  
            c = func1(i,j)  
        else:  
            c = func2(i,j)  
  
        results.append(c)  
  
final = dask.compute(results)
```

# Challenges with Large ML models

- **Large Memory Problem:**
  - Dask chunks the dataset into several pieces, so that only a fraction at a time comes into memory.
- **Large Model Problem:**
  - Dask distributes and trains sub-models onto available nodes.
  - Dask distributes batches onto multiple nodes.

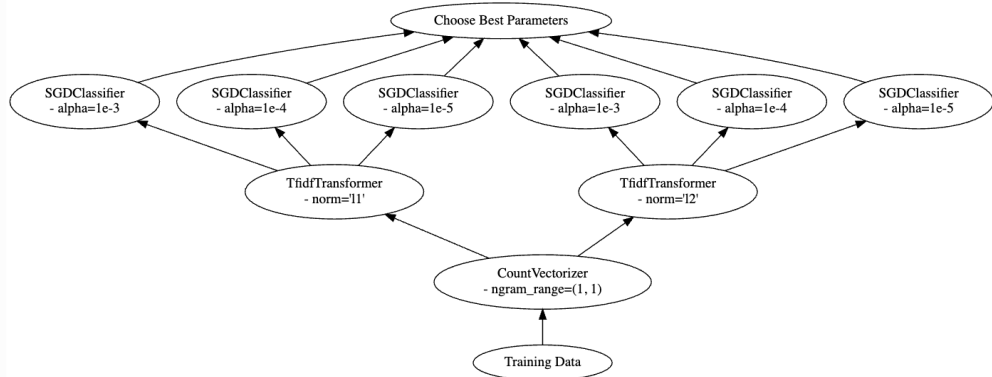
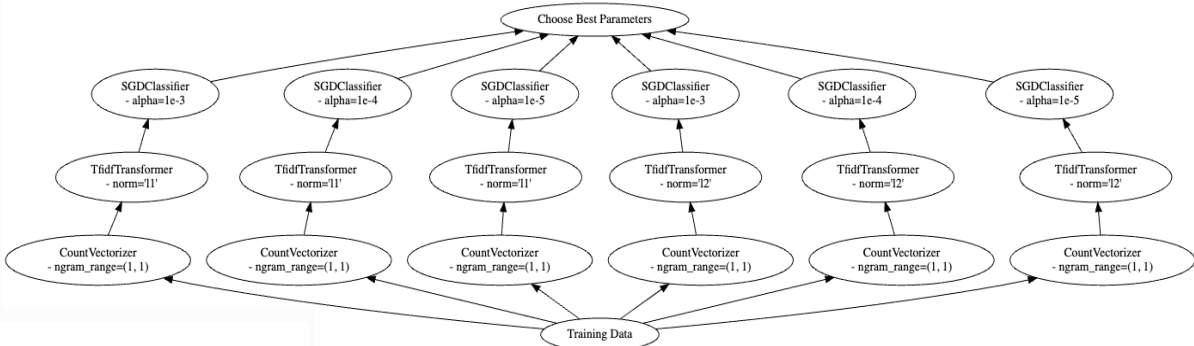


# ML jobs where Dask can help

- **HyperParameter Search**
- **Generalized Linear Models**
- **Parallel Meta-estimators**
- **Incremental Learning**
- **Text-Vectorization**
- **Automated Machine Learning (tpot)**
- **Use Dask for Batch Prediction with CNN models trained in PyTorch.**
- **Perform Hyperparameter search on tensorflow based Keras models Generalized Linear Models**

# Hyperparameter search – Impact of Dask

Dask (task-graph below) eliminates need to repeat redundant computational steps.



[`dask\_ml.model\_selection.IncrementalSearchCV\(...\)`](#)

Incrementally search for hyper-parameters on models that support `partial_fit`

[`dask\_ml.model\_selection.HyperbandSearchCV\(...\)`](#)

Find the best parameters for a particular model with an adaptive cross-validation algorithm.

[`dask\_ml.model\_selection.SuccessiveHalvingSearchCV\(...\)`](#)

Perform the successive halving algorithm [\[R424ea1a907b1-1\]](#).

[`dask\_ml.model\_selection.InverseDecaySearchCV\(...\)`](#)

Incrementally search for hyper-parameters on models that support `partial_fit`

# Notes

- **During the next session, we will go through implementation examples.**
  - Demonstration will focus on utilization of Dask primarily.
  - Understanding and Developing ML models is not the primary objective of the training.
- **For additional information, links are embedded in the Jupyter Notebooks.**
- **If there are interests in some of the models discussed, a separate training will be prepared on those subject matters.**

# Contact



## Questions and Information

<https://training.hpc.mil/>

[pet@hpc.mil](mailto:pet@hpc.mil)

[Gedion.R.Teklemariam.ctr@mail.nrl.navy.mil](mailto:Gedion.R.Teklemariam.ctr@mail.nrl.navy.mil)