



**Hewlett Packard**  
Enterprise

# **HPE Performance Software — Message Passing Interface**

## Cpuset Software Guide

**Abstract:**

This manual describes how to use cpusets. The cpuset software is part of the HPE Performance Software — Message Passing Interface software package.

Part Number: 007-5945-004

Published: August 2017

Edition: Revision 4

## **Notices**

The information contained herein is subject to change without notice. The only warranties for Hewlett Packard Enterprise products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty. Hewlett Packard Enterprise shall not be liable for technical or editorial errors or omissions contained herein.

Confidential computer software. Valid license from Hewlett Packard Enterprise required for possession, use, or copying. Consistent with FAR 12.211 and 12.212, Commercial Computer Software, Computer Software Documentation, and Technical Data for Commercial Items are licensed to the U.S. Government under vendor's standard commercial license.

Links to third-party websites take you outside the Hewlett Packard Enterprise website. Hewlett Packard Enterprise has no control over and is not responsible for information outside the Hewlett Packard Enterprise website.

## **Acknowledgements**

Linux® is the registered trademark of Linus Torvalds in the U.S. and other countries.

Red Hat® and Red Hat® Enterprise Linux are registered trademarks of Red Hat, Inc., in the United States and other countries.

UNIX® is a registered trademark and the X Window System™ is a trademark of The Open Group in the United States and other countries.

---

## Record of Revision

Version	Description
001	November 2013 Original publication.
002	May 2016 This revision supports SGI Performance Suite 1.12.
003	June 2017 This revision supports the HPE Performance Software – Message Passing Interface 1.0 release and the SGI MPI 1.14 release.
004	August 2017 Supports the HPE Performance Software – Message Passing Interface 1.0 release and the MPT 2.16 release.



---

# Contents

<b>About This Guide</b>	<b>ix</b>
HPE Performance Software Suite Publications and General MPI Information	ix
HPE Websites	xi
Documentation Feedback	xi
<b>1. Introduction</b>	<b>1</b>
About Cpusets	1
About File Names Used on Different Operating Systems	2
<b>2. Using Cpusets</b>	<b>5</b>
Cpuset Working Environment	6
Linux 2.6 Kernel Support for Cpusets	6
Cpuset Facility Capabilities	7
Initializing Cpusets	8
Verifying Whether Cpusets are Installed and Whether the Correct Kernel is Installed	8
Fine-grained Control within Cpusets	9
Cpuset Interaction with Other Placement Mechanisms	9
Cpusets and Thread Placement	11
Safe Job Migration and Cpusets	12
Application Performance	13
Cpuset File System Directories	14
Exclusive Cpusets	19
Notify on Release Flag	19
Memory Pressure of a Cpuset	20

Memory Spread . . . . .	22
Memory Migration . . . . .	23
Mask Format . . . . .	24
List Format . . . . .	25
Cpuset Permissions . . . . .	25
CPU Scheduling and Memory Allocation for Cpusets . . . . .	26
Linux Kernel CPU and Memory Placement Settings . . . . .	26
Manipulating Cpusets . . . . .	27
Using Cpusets at the Shell Prompt . . . . .	27
Cpuset Command Line Utility . . . . .	29
Action Options . . . . .	31
Modifier Options . . . . .	33
Help Options . . . . .	34
Boot Cpuset . . . . .	34
Creating a Bootcpuset . . . . .	34
bootcpuset.conf File . . . . .	35
Configuring a User Cpuset for Interactive Sessions . . . . .	36
Cpuset Text Format . . . . .	37
Modifying the CPUs in a Cpuset and Kernel Processing . . . . .	38
Using Cpusets with Hyper-Threads . . . . .	39
Cpuset Programming Model . . . . .	41
System Error Messages . . . . .	42
<b>Appendix A. Cpuset Library Functions . . . . .</b>	<b>45</b>
Basic Cpuset Library Functions . . . . .	45
Extensible Application Programming Interface . . . . .	46
Advanced Cpuset Library Functions . . . . .	47

<b>Index</b>	<b>55</b>
--------------	-----------





---

## About This Guide

This manual describes how to use cpusets. The cpuset software is part of the HPE Performance Software — Message Passing Interface software package.

You can use cpusets to restrict the number of processors and the amount of memory that a process or a set of processes can use.

## HPE Performance Software Suite Publications and General MPI Information

The HPE Performance Software Suite Publications are as follows:

- The *HPE Performance Software - Message Passing Interface User Guide* describes how to use the HPE Performance Software — Message Passing Interface software.
- The *HPE Performance Software - Message Passing Interface MPInside Reference Guide* describes the MPI profiling tool called MPInside.
- The *HPE Performance Software - Message Passing Interface Cpuset Software Guide* describes how to use cpusets.
- MPInside(3)

This man page lists all the MPInside environment variables that HPE supports.

- MPInside-exp(3)

This man page lists MPInside environment variables that HPE supports on an experimental basis. Use of these environment variables can generate unexpected results. The HPE Performance MPI documentation uses some of these experimental variables in examples and procedures.

- The HPE Performance Software — Message Passing Interface release notes contain information about the specific software packages included in each product. The release notes are available in the following locations:

- The HPE Performance MPI release notes are posted to the following SGI website:

[https://support1-sgi.custhelp.com/app/answers/detail/a\\_id/6093](https://support1-sgi.custhelp.com/app/answers/detail/a_id/6093)

---

**Note:** To access documentation through the SGI customer portal, first log in, and then navigate to the link. You can log in at the following website:

<https://support.sgi.com/login>

---

- On the product media. The release notes reside in a text file in the `/docs` directory on the product media.
- On the system. After installation, the release notes and other product documentation reside in the `/opt/hpe/hpc/mpt/mpt-version_number/doc/README.relnotes` directory.

For example, you can use the following `rpm(8)` command to retrieve the location of the Array Services release notes:

```
# rpm -qi sgi-arraysvcs
/usr/share/doc/packages/sgi-arraysvcs/README.relnotes
```

Use a text editor or other command to display the file that the `rpm(8)` command returns.

For information about the MPI standard, see the following:

- The Message Passing Interface Forum’s website, which is as follows:  
<http://www.mpi-forum.org/>
- *Using MPI — 2nd Edition: Portable Parallel Programming with the Message Passing Interface (Scientific and Engineering Computation)*, by Gropp, Lusk, and Skjellum. ISBN-13: 978-0262571326.
- The University of Tennessee technical report. See reference [24] from *Using MPI: Portable Parallel Programming with the Message-Passing Interface*, by Gropp, Lusk, and Skjellum. ISBN-13: 978-0262571043.
- Journal articles in the following publications:
  - *International Journal of Supercomputer Applications*, volume 8, number 3/4, 1994
  - *International Journal of Supercomputer Applications*, volume 12, number 1/4, pages 1 to 299, 1998

## HPE Websites

The following HPE websites might interest you:

- Hewlett Packard Enterprise Information Library:  
[www.hpe.com/info/EIL](http://www.hpe.com/info/EIL)
- Single Point of Connectivity Knowledge (SPOCK) Storage compatibility matrix:  
[www.hpe.com/storage/spock](http://www.hpe.com/storage/spock)
- Storage white papers and analyst reports:  
[www.hpe.com/storage/whitepapers](http://www.hpe.com/storage/whitepapers)

## Documentation Feedback

Hewlett Packard Enterprise is committed to providing documentation that meets your needs. To help us improve the documentation, send any errors, suggestions, or comments to Documentation Feedback ([docsfeedback@hpe.com](mailto:docsfeedback@hpe.com)). When submitting your feedback, include the document title, part number, edition, and publication date located on the front cover of the document. For online help content, include the product name, product version, help edition, and publication date located on the legal notices page.



## Introduction

This chapter contains the following topics:

- "About Cpusets" on page 1
- "About File Names Used on Different Operating Systems" on page 2

### About Cpusets

A *cpuset* defines a list of CPUs and memory nodes. A process contained in a *cpuset* can only execute on the CPUs in that *cpuset* and can only allocate memory on the memory nodes in that *cpuset*. Essentially, *cpusets* provide you with a CPU and memory containers, or soft partitions, within which you can run sets of related tasks.

The *cpuset* facility allows a system administrator or a workload manager, such as PBS Pro or MOAB/Torque, to restrict the number of processor and memory resources that a process or set of processes can use.

If you use *cpusets* on an SGI UV system, for example, you can improve memory locality, memory access times, and an application's performance and runtime repeatability. Restraining all other jobs from using any of the CPUs or memory resources assigned to a critical job minimizes interference from other jobs on the system. For example, Message Passing Interface (MPI) jobs frequently consist of a number of threads that communicate using message passing interfaces. All threads need to be executing at the same time. If a single thread loses a CPU, all threads stop making forward progress and spin at a synchronization or communication function.

The following are some of the advantages to using *cpusets*:

- *Cpusets* can eliminate the need for a gang scheduler, provide isolation of one such job from other tasks on a system, and facilitate providing equal resources to each thread in a job. This results in both optimum and repeatable performance.

In addition to their traditional use to control the placement of jobs on the CPUs and memory nodes of a system, *cpusets* also provide a convenient mechanism to control the use of Hyper-Threading Technology.

- *Cpusets* are represented in a hierarchical virtual file system. *Cpusets* can be nested and they have file-like permissions.

- The `sched_setaffinity`, `mbind`, and `set_mempolicy` system calls allow you to specify the CPU and memory placement for individual tasks. On smaller or limited-use systems, these calls might be sufficient.
- The kernel cpuset facility provides additional support for system-wide management of CPU and memory resources by related sets of tasks. It provides a hierarchical structure to the resources, with filesystem-like namespace and permissions, and support for guaranteed exclusive use of resources.
- You can have a *boot* cpuset running the traditional daemon and server tasks and a second cpuset to hold interactive `telnet`, `rlogin` and/or secure shell (SSH) user sessions called the *user* cpuset.

Creating a user cpuset provides additional isolation between interactive user login sessions and essential system tasks. For example, a user process in the user cpuset that consumes excessive CPU, system file buffer cache, or memory resources does not seriously affect essential system services in the boot cpuset. For more information, see "Configuring a User Cpuset for Interactive Sessions" on page 36.

## About File Names Used on Different Operating Systems

This documentation has been updated for the Red Hat Enterprise Linux (RHEL) 7 operating system and the SLES 12 operating system. If you use these operating systems, the `/dev/cpusets` directory includes many file names that begin with a `cpuset.` prefix. The following files are among those that reside in the `/dev/cpusets` directory:

```
cpuset.cpu_exclusive
cpuset.cpus
cpuset.mem_exclusive
cpuset.mem_hardwall
cpuset.memory_migrate
cpuset.memory_pressure
cpuset.memory_pressure_enabled
cpuset.memory_spread_slab
cpuset.memory_spread_page
cpuset.mems
cpuset.sched_load_balance
cpuset.sched_relax_domain_level
```

If you use an older operating system, none of the files in the `/dev/cpusets` directory include the `cpuset.` prefix. The names for these same files are as follows:

```
cpu_exclusive
cpus
mem_exclusive
mem_hardwall
memory_migrate
memory_pressure
memory_pressure_enabled
memory_spread_slab
memory_spread_page
mems
sched_load_balance
sched_relax_domain_level
```





## Using Cpusets

This chapter contains the following topics:

- "Cpuset Working Environment" on page 6
- "Cpuset File System Directories" on page 14
- "Exclusive Cpusets" on page 19
- "Notify on Release Flag" on page 19
- "Memory Pressure of a Cpuset" on page 20
- "Memory Spread" on page 22
- "Memory Migration" on page 23
- "Mask Format" on page 24
- "List Format" on page 25
- "Cpuset Permissions" on page 25
- "CPU Scheduling and Memory Allocation for Cpusets" on page 26
- "Using Cpusets at the Shell Prompt" on page 27
- "Cpuset Command Line Utility" on page 29
- "Boot Cpuset" on page 34
- "Configuring a User Cpuset for Interactive Sessions" on page 36
- "Cpuset Text Format" on page 37
- "Modifying the CPUs in a Cpuset and Kernel Processing" on page 38
- "Using Cpusets with Hyper-Threads" on page 39
- "Cpuset Programming Model" on page 41
- "System Error Messages" on page 42

## Cpuset Working Environment

This section covers the following topics:

- "Linux 2.6 Kernel Support for Cpusets" on page 6
- "Cpuset Facility Capabilities" on page 7
- "Initializing Cpusets" on page 8
- "Verifying Whether Cpusets are Installed and Whether the Correct Kernel is Installed" on page 8
- "Fine-grained Control within Cpusets" on page 9
- "Cpuset Interaction with Other Placement Mechanisms" on page 9
- "Cpusets and Thread Placement" on page 11
- "Safe Job Migration and Cpusets" on page 12

## Linux 2.6 Kernel Support for Cpusets

The Linux 2.6 kernel provides the following support for cpusets:

- Each task has a link to a cpuset structure that specifies the CPUs and memory nodes available for its use.
- Hooks in the `sched_setaffinity` system call, used for CPU placement, and in the `mbind` system call, used for memory placement, ensure that any requested CPU or memory node is available in that task's cpuset.
- All tasks sharing the same placement constraints reference the same cpuset.
- Kernel cpusets are arranged in a hierarchical virtual file system, reflecting the possible nesting of soft partitions.
- The kernel task scheduler is constrained to only schedule a task on the CPUs in that task's cpuset.
- The kernel memory allocation mechanism is constrained to only allocate physical memory to a task from the memory nodes in that task's cpuset.
- The kernel memory allocation mechanism provides an economical, per-cpuset metric of the aggregate memory pressure of the tasks in a cpuset. *Memory pressure* is defined as the frequency of requests for a free memory page that is not easily

satisfied by an available free page. For more information, see "Memory Pressure of a Cpuset" on page 20.

- The kernel memory allocation mechanism enables you to spread memory pages used for file I/O (the kernel page cache) and associated kernel data structures for file inodes and directories evenly across all the memory nodes in a cpuset. Otherwise, the memory pages are allocated on the memory node upon which the task first accessed the memory page.
- You can control the memory migration facility in the kernel using per-cpuset files. When the memory nodes allowed to a task by cpusets change, any memory pages no longer allowed on that node may be migrated to nodes now allowed. For more information, see "Safe Job Migration and Cpusets" on page 12.

## Cpuset Facility Capabilities

A cpuset constrains the jobs (set of related tasks) running in it to a subset of the system's memory and CPUs. The cpuset facility allows you and your system service software to do the following:

- Create and delete named cpusets.
- Decide which CPUs and memory nodes are available to a cpuset.
- Attach a task to a particular cpuset.
- Identify all tasks sharing the same cpuset.
- Exclude any other cpuset from overlapping a given cpuset, thereby, giving the tasks running in that cpuset exclusive use of those CPUs and memory nodes.
- Perform bulk operations on all tasks associated with a cpuset, such as varying the resources available to that cpuset or hibernating those tasks in temporary favor of some other job.
- Perform sub-partitioning of system resources using hierarchical permissions and resource management.

### Initializing Cpusets

The kernel, at system boot time, initializes one cpuset, the root cpuset, containing the entire system's CPUs and memory nodes. Subsequent user space operations can create additional cpusets.

Mounting the cpuset virtual file system (VFS) at `/dev/cpuset` exposes the kernel mechanism to user space. This VFS allows for nested resource allocations and the associated hierarchical permission model.

You can initialize and perform other cpuset operations, using any of the following mechanisms:

- You can create, change, or query cpusets by using shell commands on `/dev/cpuset`, such as `echo(1)`, `cat(1)`, `mkdir(1)`, or `ls(1)` as described in "Using Cpusets at the Shell Prompt" on page 27.
- You can use the `cpuset(1)` command line utility to create or destroy cpusets or to retrieve information about existing cpusets and to attach processes to existing cpusets as described in "Cpuset Command Line Utility" on page 29.
- You can use the `libcpuset` C programming application programming interface (API) functions to query or change them from within your application as described in Appendix A, "Cpuset Library Functions" on page 45. You can find information about `libcpuset` at </usr/share/doc/packages/libcpuset/libcpuset.html>.

### Verifying Whether Cpusets are Installed and Whether the Correct Kernel is Installed

The following procedure explains how to use commands to determine whether cpusets are installed on your system.

**Procedure 2-1** To determine whether cpusets and the correct kernel is installed

1. Use the `grep(1)` command to search the `/proc/filesystems` for cpusets, as follows:

```
% grep cpuset /proc/filesystems
nodev    cpuset
```

2. Change to the `/dev/cpuset` directory and list the content of the directory.

For example:

```
% cd /dev/cpuset
Directory: /dev/cpuset

% ls
cpuset.cpu_exclusive  cpuset.cpus  cpuset.mem_exclusive  cpuset.mems
notify_on_release    pagecache_list  pagecache_local  slabcache_local
tasks
```

The preceding example shows that the `tasks` file exists, so in this case, cpusets are installed.

If the `/dev/cpuset/tasks` file is not present, it means that the cpuset file system is not mounted. Usually, it is automatically mounted when the system is booted. As root, you can use the following command to mount the cpuset file system:

```
% mount -t cpuset cpuset /dev/cpuset
```

3. Use the `uname(1)` command to retrieve kernel information, and make sure you are running the Linux 2.6.x `sn2` kernel:

```
% uname -r -s
Linux 2.6.16.14-6-default
```

## Fine-grained Control within Cpusets

Within a single cpuset, use facilities such as `taskset(1)`, `dplace(1)`, `first-touch` memory placement, `pthread_setaffinity_np`, and `mbind` to manage processor and memory placement to a more fine-grained level.

The user-level bitmask library supports convenient manipulation of multiword bitmasks useful for CPUs and memory nodes. This bitmask library is required by and designed to work with the cpuset library. You can find information about the bitmask library on your system at </usr/share/doc/packages/libbitmask/libbitmask.html>.

## Cpuset Interaction with Other Placement Mechanisms

The Linux 2.6 kernel supports the following additional processor and memory placement mechanisms:

- The `sched_setaffinity(2)` and `sched_getaffinity(2)` system calls set and get the CPU affinity mask of a process. This determines the set of CPUs on which a process is eligible to run. The `taskset(1)` command provides a command line utility for manipulating the CPU affinity mask of a process using these system calls. For more information, see the appropriate man page.
- The `set_mempolicy` system call sets the NUMA memory policy of the current process to *policy*. A NUMA machine has different memory controllers with different distances to specific CPUs. The memory *policy* defines in which node memory is allocated for the process.

The `get_mempolicy(2)` system retrieves the NUMA policy of the calling process or of a memory address, depending on the setting of *flags*. The `numactl(8)` command provides a command line utility for manipulating the NUMA memory policy of a process using these system calls.

- The `mbind(2)` system call sets the NUMA memory policy for the pages in a specific range of a task's virtual address space.

Cpusets are designed to interact cleanly with other placement mechanisms. For example, a workload manager can use cpusets to control the CPU and memory placement of various jobs; while within each job, these other kernel mechanisms are used to manage placement in more detail. It is possible for a workload manager to change a job's cpuset placement while preserving the internal CPU affinity and NUMA memory placement policy, without requiring any special coding or awareness by the affected job.

Most jobs initialize their placement early in their time slot, and jobs are rarely migrated until they have been running for a while. As long as a workload manager does **not** try to migrate a job at the same time as it is adjusting its own CPU or memory placement, there is little risk of interaction between cpusets and other kernel placement mechanisms.

The CPU and memory node placement constraints imposed by cpusets always override those of these other mechanisms.

Calls to the `sched_setaffinity(2)` system call automatically mask off CPUs that are not allowed by the affected task's cpuset. If a request results in all the CPUs being masked off, the call fails with `errno` set to `EINVAL`. If some of the requested CPUs are allowed by the task's cpuset, the call proceeds as if only the allowed CPUs were requested. The disallowed CPUs are silently ignored. If a task is moved to a different cpuset, or if the CPUs of a cpuset are changed, the CPU affinity of the affected task or tasks is lost. If a workload manager needs to preserve the CPU affinity of the tasks in

a job that is being moved, it should use the `sched_setaffinity(2)` and `sched_getaffinity(2)` calls to save and restore each affected task's CPU affinity across the move, relative to the cpuset. The `cpu_set_t` mask data type supported by the C library for use with the CPU affinity calls is different from the `libbitmask` bitmasks used by `libcpuset`, so some coding will be required to convert between the two, in order to calculate and preserve cpuset relative CPU affinity.

Similar to CPU affinity, calls to modify a task's NUMA memory policy silently mask off requested memory nodes outside the task's allowed cpuset. The calls to modify a task's NUMA memory policy fail if that results in an empty set of memory nodes. Unlike CPU affinity, the NUMA memory policy system calls do not support one task querying or modifying another task's policy. The kernel automatically handles preserving cpuset relative to NUMA memory policy when either a task is attached to a different cpuset or when a cpuset's `mems` value setting is changed. If the old and new `mems` value sets have the same size, the cpuset relative offset of affected NUMA memory policies is preserved. If the new `mems` value is smaller, the old `mems`-value-relative offsets are folded onto the new `mems` value, modulo the size of the new `mems`. If the new `mems` value is larger, then just the first N nodes are used, where N is the size of the old `mems` value.

## Cpusets and Thread Placement

If your job uses the placement mechanisms described in "Cpuset Interaction with Other Placement Mechanisms" on page 9 and operates under the control of a workload manager, you **cannot** guarantee that a migration will preserve placement done using the mechanisms. These placement mechanisms use system wide numbering of CPUs and memory nodes, not cpuset relative numbering, and the job might be migrated without its knowledge while it is trying to adjust its placement. That is, between the point where an application computes the CPU or memory node on which it wants to place a thread and the point at which it issues the `sched_setaffinity(2)`, `mbind(2)`, or `set_mempolicy(2)` call to direct such a placement, the thread might be migrated to a different cpuset, or the thread's cpuset might change to different CPUs or memory nodes, invalidating the CPU or memory node number it just computed.

The `libcpuset` library provides the following mechanisms to support cpuset-relative thread placement that is robust even if the job is being migrated using a batch scheduler.

If your job needs to pin a thread to a single CPU, you can use the convenient `cpuset_pin` function. This is the most common case. For more information on `cpuset_pin`, see the `libcpuset(3)` man page.

If your job needs to implement some other variation of placement, such as to specific memory nodes, or to more than one CPU, use one of the following functions to safely guard such code from placement changes caused by job migration:

- `cpuset_get_placement`
- `cpuset_equal_placement`
- `cpuset_free_placement`

For information about the placement functions, see `libcpuset_placement(3)`.

### Safe Job Migration and Cpusets

Jobs that make use of cpuset-aware thread pinning can be safely migrated to a different cpuset or can have the CPUs or memory nodes of the cpuset safely changed without destroying the per-thread placement done within the job.

When you migrate a job from a larger cpuset (more CPUs or nodes) to a smaller cpuset, the result is a loss of placement information. You cannot move that cpuset back to a larger cpuset to recover the placement information. To avoid the loss of CPU affinity, use `sched_getaffinity(2)` and `sched_setaffinity(2)` to save and restore the placement (affinity) across such a pair of moves. This loss of NUMA memory placement information cannot be avoided because one task (the one doing the migration) can neither save nor restore the NUMA memory placement policy of another. So if a workload manager wants to migrate jobs without causing them to lose their `mbind(2)` or `set_mempolicy(2)` placement, it should only migrate to cpusets with at least as many memory nodes as the original cpuset.

For information about cpuset-aware thread pinning, see the following:

"Cpusets and Thread Placement" on page 11

**Procedure 2-2** To migrate a job between cpusets

1. Suspend the tasks in the job by sending their process group a `SIGSTOP` signal.
2. Use the `cpuset_init_pidlist` function and related `pidlist` functions to determine the list of tasks in the job.
3. Use `sched_getaffinity(2)` to query the CPU affinity of each task in the job.

The `sched_getaffinity(2)` and `sched_setaffinity(2)` C library calls are limited by C library internals to systems with 1024 CPUs or less. To write code that works on larger systems, use the `syscall(2)` indirect system call wrapper to



directly invoke the underlying system call, bypassing the C library API for these calls.

The suspend and resume operations are required in order to keep tasks in the job from changing their per-thread CPU placement between this step and the following step:

Procedure 2-2, step 6 on page 13

4. Create a new cpuset, under a temporary name, with the new desired CPU and memory placement.

The kernel automatically migrates the per-thread memory node placement during this step. This is necessary because there is no way for one task to modify the NUMA memory placement policy of another task. The kernel does not automatically migrate the per-thread CPU placement because this can be handled by the user-level process doing the migration.

5. Invoke `cpuset_migrate_all` function to move the job's tasks from the old cpuset to the new cpuset.
6. Use `cpuset_delete` to delete the old cpuset.
7. Use `rename(2)` on the `/dev/cpuset` based path of the new temporary cpuset to rename that cpuset to the old cpuset name.
8. Convert the results of the previous `sched_getaffinity(2)` calls to the new cpuset placement, preserving cpuset relative offset by using the `cpuset_c_rel_to_sys_cpu` and related functions.
9. Use `sched_setaffinity(2)` to reestablish the per-task CPU binding of each thread in the job.
10. Resume the tasks in the job by sending their process group a `SIGCONT` signal.

## Application Performance

The `cpuset(7)` man page lists several settings that you need to monitor when you run applications on the following systems:

- HPE SGI 8600 systems
- HPE Apollo 6500 systems
- HPE Apollo 4000 systems

- HPE Apollo 2000 systems
- HPE ProLiant systems
- HPE Integrity MC990 X systems
- SGI UV systems

The settings to monitor are as follows:

- `mem_exclusive`

When set (1), the cpuset has exclusive use of its memory nodes. No sibling or cousin can overlap. Also, when set, the cpuset is a hardwall cpuset.

By default, this is off (0).

- `mem_spreadpage`

When set (1), pages in the kernel page cache (file-system buffers) are uniformly spread across the cpuset.

By default, this is off (0) in the top cpuset and is inherited from the parent cpuset in newly created cpusets.

- `mem_hardwall`

When set (1), the cpuset is a hardwall cpuset. Unlike `mem_exclusive`, there is no constraint on whether cpusets marked `mem_hardwall` can have overlapping memory nodes with sibling or cousin cpusets.

By default, this is off (0). Newly created cpusets also initially default this to off (0).

For more information, see the `cpuset(7)` man page.

## Cpuset File System Directories

Cpusets are named, nested sets of CPUs and memory nodes. Each cpuset is represented by a directory in the cpuset virtual file system, typically mounted at `/dev/cpuset`.

The state of each cpuset is represented by small text files in the directory for the cpuset. To read and write these files using traditional shell utilities, use `cat(1)` or `echo(1)`. To access these files from programming languages, use `open(2)`, `read(2)`, `write(2)`, and `close(2)` from the C programming library.

To view the files in a cpuset that can be either read or written, type the following commands:

```
% cd /dev/cpuset
% ls
```

This topic describes the files in the cpuset directory. Many of these files accept a flag setting that controls cpuset actions. To change the action, write an ASCII number 1 or 0 (with an optional trailing newline) into the file. A 1 sets the flag. A 0 clears the flag.

The following list describes the files in the cpuset directory:

File	Description
tasks	<p>Lists the process IDs (PIDs) of tasks in the cpuset. The list is formatted as a series of ASCII decimal numbers, each followed by a newline. A task can be added to a cpuset, which removes it from the cpuset that previously contained it, by writing its PID to that cpuset's tasks file, with or without a trailing newline.</p> <p>Write only one PID to the tasks file at a time. If a string is written that contains more than one PID, all but the first are ignored.</p>
notify_on_release	<p>When set (1), the system invokes the <code>/sbin/cpuset_release_agent</code> binary with the name <code>(/dev/cpuset relative path)</code> of the cpuset in <code>argv[1]</code> when the last user of it (task or child cpuset) goes away. This supports automatic cleanup of abandoned cpusets.</p> <p>When not set (0), automatic cleanup of abandoned cpusets does not occur.</p> <p>For more information, see the following:</p> <p>"Notify on Release Flag" on page 19</p>
cpuset.cpus	<p>Lists the CPUs that tasks in the cpuset are allowed to use. To change the CPUs allowed to a cpuset, write a new list to its <code>cpuset.cpus</code></p>

file. Note, however, such a change does not take effect until the PIDs of the tasks in the cpuset are rewritten to the cpuset's `tasks` file.

For a description of the format of the `cpuset.cpus` file, see "List Format" on page 25.

`cpuset.cpu_exclusive`

When set (1), the system automatically defines scheduler domains. The kernel performs automatic load balancing of active threads on available CPUs more rapidly within a scheduler domain than it does across scheduler domains.

When not set (0), which is the default, newly created cpusets initially default this flag to off.

`cpuset.mems`

Lists memory nodes that tasks in the cpuset are allowed to use.

For a description of the format of the `mems` file, see "List Format" on page 25.

`cpuset.mem_exclusive`

When set (1), the `cpuset.mem_exclusive` flag automatically defines constraints for kernel internal memory allocations. Allocations of user space memory pages are strictly confined by the allocating task's cpuset. Allocations of kernel internal pages are only confined by the nearest enclosing cpuset that is marked `cpuset.mem_exclusive`.

When not set (0), newly created cpusets initially default this flag to off.

`cpuset.memory_migrate`

When set (1), memory migration is enabled.

For more information, see "Memory Migration" on page 23.

`cpuset.memory_pressure`

This file contains a measure of how much memory pressure the tasks in this cpuset are causing. This is a read-only file.

Always has value zero (0) unless `cpuset.memory_pressure_enabled` is enabled in the top cpuset. The `cpuset.memory_pressure` mechanism makes it easy to detect when the job in a cpuset is running short of memory and needing to page memory out to swap.

For more information, see "Memory Pressure of a Cpuset" on page 20.

`cpuset.memory_pressure_enabled`

This file is only present in the root cpuset, typically at `/dev/cpuset`.

When set (1), `cpuset.memory_pressure` calculations are enabled for all cpusets on the system.

For more information, see "Memory Pressure of a Cpuset" on page 20.

`cpuset.memory_spread_page`

When set (1), the kernel page cache (file system buffers) are uniformly spread across the cpuset.

For more information, see "Memory Spread" on page 22.

`cpuset.memory_spread_slab`

When set (1), the kernel slab caches for file I/O (directory and inode structures) are uniformly spread across the cpuset.

For more information, see "Memory Spread" on page 22.

The `/proc/pid/cpuset` file exists in the `/proc` file system. For each task (PID), the file lists its cpuset path, relative to the root of the cpuset file system. This is a read-only file.

Within `/proc/pid/status`, the kernel scheduler and memory allocation mechanism use two control fields to constrain scheduling and memory allocation to the allowed CPUs. The control fields are as follows:

<code>Cpus_allowed</code>	A bit vector of CPUs on which this task can be scheduled
<code>Mems_allowed</code>	A bit vector of memory nodes on which this task can obtain memory

There are several reasons why a task's `Cpus_allowed` and `Mems_allowed` values might differ from the values in the `cpus` file and the `mems` file. For example:

- A task might use the `sched_setaffinity`, `mbind`, or `set_mempolicy` functions to restrain its placement to less than its cpuset.
- Various temporary changes to `cpus_allowed` values are done by kernel internal code.
- Attaching a task to a cpuset does not change its `mems_allowed` value until the next time that task needs kernel memory.
- Changing a cpuset's `cpus` value does not change the `Cpus_allowed` value of the tasks attached to the cpuset until those tasks reattach to that cpuset. This practice avoids a hook in the hotpath scheduler code in the kernel.

User space action is required to update a task's `Cpus_allowed` values after changing its cpuset. Use the `cpuset_reattach` routine to perform this update after changing the CPUs allowed to a cpuset.

- If the hotplug mechanism is used to remove all the CPUs or all the memory nodes in a cpuset, the tasks attached to that cpuset are updated to have their `Cpus_allowed` or `Mems_allowed` values altered to the CPUs or memory nodes of the closest ancestor to that cpuset that is not empty.

The confines of a cpuset can be violated after a hotplug removal that empties a cpuset. To avoid this situation, the system's cpuset configuration is updated. The update reflects the new hardware configuration and does not define a cpuset that has no CPUs still online or no memory nodes still online. The kernel prefers to misplace a task rather than to starve a task of essential compute resources.

There is one other condition under which the confines of a cpuset can be violated. A few kernel-critical internal memory allocation requests, marked `GFP_ATOMIC`, must be satisfied immediately. If one of these allocations fails, the kernel might drop a request or malfunction. If such a request cannot be satisfied within the current task's cpuset, the kernel relaxes the cpuset and looks for memory anywhere it can find it. It is better to violate the cpuset than to stress the kernel operation.

To create new cpusets, use the `mkdir` command at the shell (see "Using Cpusets at the Shell Prompt" on page 27) or use the C programming language (see Appendix A, "Cpuset Library Functions" on page 45). To remove old cpusets, use the `rmdir(1)` command.

The root cpuset includes all possible CPUs and memory nodes in the system. A cpuset can be exclusive to specific CPUs or to specific memory nodes only if the cpuset's parent is similarly exclusive.

Each task has a pointer to a cpuset. Multiple tasks can reference the same cpuset. Requests by a task are filtered through that task's cpuset. The system filters out any CPUs or memory nodes that are not in that cpuset. Task requests can use the `sched_setaffinity(2)` system call to include CPUs in its CPU affinity mask and can use the `mbind(2)` and `set_mempolicy(2)` system calls to include memory nodes in its memory policy. The scheduler does not schedule a task on a CPU that is not allowed in its `cpus_allowed` vector. The kernel page allocator does not allocate a page on a node that is not allowed in the requesting task's `mems_allowed` vector.

## Exclusive Cpusesets

If a cpuset is marked `cpu_exclusive` or `mem_exclusive`, only a direct ancestor cpuset or direct descendant cpuset can share any of the same CPUs or memory nodes. The effect is as follows:

- A cpuset that is `cpu_exclusive` is associated with a scheduler (`sched`) domain. The `sched` domain consists of all CPUs in the current cpuset that are not part of any exclusive child cpusesets. These boundaries ensure that the scheduler load balancing code balances only against the CPUs that are in the `sched` domain as described in "Cpuset File System Directories" on page 14 and not all of the CPUs in the system. This practice removes any overhead incurred when the load balancing code tries to pull tasks out of the `cpu_exclusive` cpuset only to be prevented by the `cpus_allowed` mask of the task.
- A cpuset that is `mem_exclusive` restricts kernel allocations for page, buffer, and other data that is commonly shared by the kernel across multiple users. All cpusesets, whether `mem_exclusive` or not, restrict allocations of memory for user space. This practice enables a system configuration in which several independent jobs share common kernel data, such as file system pages, while isolating the user allocation of each job to its own cpuset. To do this, construct a large `mem_exclusive` cpuset to hold all the jobs, and construct child, non-`mem_exclusive` cpusesets for each individual job. Only a small amount of typical kernel memory, such as requests from interrupt handlers, is allowed to be taken outside even a `mem_exclusive` cpuset.

## Notify on Release Flag

When set (1), the `notify_on_release` flag removes abandoned cpusesets automatically. When the last task in the cpuset leaves, exits, or attaches to some other

cpuset, and the last child cpuset of that cpuset is removed, the kernel runs the `/sbin/cpuset_release_agent` command. The kernel uses the path name of the abandoned cpuset, relative to the mount point of the cpuset file system, as the argument to the `/sbin/cpuset_release_agent` command.

The typical content of the `/sbin/cpuset_release_agent` command is the following simple shell script:

```
#####  
# NOTE: The example shown is for illustrative purposes only and #  
# has not been evaluated for use in a production environment.  #  
#####  
#!/bin/sh  
rmdir /dev/cpuset/$1
```



---

**Caution:** The preceding script example is for illustrative purposes only and has not been evaluated for use in a production environment.

---

By default, the `notify_on_release` flag is off (0). Newly created cpusets inherit their `notify_on_release` flag setting from their parent cpuset. As with other flag values, this flag can be changed by writing an ASCII number 0 or 1 (with optional trailing newline) into the file, to clear or set the flag, respectively.

## Memory Pressure of a Cpuset

The `memory_pressure` of a cpuset provides a simple per-cpuset metric of the rate that the tasks in a cpuset are attempting to free up in use memory on the nodes of the cpuset to satisfy additional memory requests. This enables workload managers, monitoring jobs running in dedicated cpusets, to efficiently detect what level of memory pressure that job is causing.

This is useful in the following situations:

- Tightly managed systems running a wide mix of submitted jobs that may choose to terminate or re-prioritize jobs trying to use more memory than allowed on the nodes to which they are assigned.
- Tightly coupled, long running, massively parallel, scientific computing jobs that will dramatically fail to meet required performance goals if they start to use more memory than allowed.



This mechanism provides a very economical way for the workload manager to monitor a cpuset for signs of memory pressure. It is up to the workload manager or other user code to decide when to take action to alleviate memory pressures.

If the `cpuset.memory_pressure_enabled` flag in the top cpuset is (0), that is, it is **not** set, the kernel does not compute this filter and the per-cpuset files `cpuset.memory_pressure` contain the value zero (0).

If the `cpuset.memory_pressure_enabled` flag in the top cpuset is set (1), the kernel computes this filter for each cpuset in the system, and the `cpuset.memory_pressure` file for each cpuset reflects the recent rate of such low memory page allocation attempts by tasks in said cpuset.

Reading the `cpuset.memory_pressure` file of a cpuset is very efficient. This mechanism allows batch schedulers to poll these files and detect jobs that are causing memory stress. They can then take action to avoid impacting the rest of the system with a job that is trying to aggressively exceed its allowed memory.

---

**Note:** Unless enabled by setting `cpuset.memory_pressure_enabled` in the top cpuset, `cpuset.memory_pressure` is not computed for any cpuset and always reads a value of zero.

---

A running average per cpuset has the following advantages:

- The system load imposed by a batch scheduler monitoring this metric is sharply reduced on large systems because this meter is per-cpuset, rather than per-task or memory region and this avoids a scan of the system-wide task list on each set of queries.
- A batch scheduler can detect memory pressure with a single read, instead of having to read and accumulate results for a period of time because this meter is a running average, rather than an accumulating counter.
- A batch scheduler can obtain the key information, memory pressure in a cpuset, with a single read, rather than having to query and accumulate results over all the (dynamically changing) set of tasks in the cpuset because this meter is per-cpuset rather than per-task or memory region.

A simple, per-cpuset digital filter is kept within the kernel and updated by any task attached to that cpuset if it enters the synchronous (direct) page reclaim code.

The per-cpuset `cpuset.memory_pressure` file provides an integer number representing the recent (half-life of 10 seconds) rate of direct page reclaims caused by the tasks in the cpuset in units of reclaims attempted per second, times 1000.

The kernel computes this value using a single-pole, low-pass recursive digital filter coded with 32-bit integer arithmetic. The value decays at an exponential rate.

Given the simple 32-bit integer arithmetic used in the kernel to compute this value, this meter works best for reporting page reclaim rates between one per millisecond (msec) and one per 32 (approximate) seconds. At constant rates faster than one per msec, it reaches maximum at values just under 1,000,000. At constant rates between one per msec and one per second, it stabilizes to a value  $N*1000$ , where  $N$  is the rate of events per second. At constant rates between one per second and one per 32 seconds, it is choppy, moving up on the seconds that have an event, and then decaying until the next event. At rates slower than about one in 32 seconds, it decays all the way back to zero between each event.

## Memory Spread

There are two Boolean flag files per cpuset that control where the kernel allocates pages for the file system buffers and related in kernel data structures. They are called `cpuset.memory_spread_page` and `cpuset.memory_spread_slab`.

If the per-cpuset, `memory_spread_page` flag is set, the kernel spreads the file system buffers (page cache) evenly over all the nodes that the faulting task is allowed to use, instead of preferring to put those pages on the node where the task is running.

If the per-cpuset, `cpuset.memory_spread_slab` flag is set, the kernel spreads some file system related slab caches, such as for inodes and directory entries, evenly over all the nodes that the faulting task is allowed to use, instead of preferring to put those pages on the node where the task is running.

The setting of these flags does not affect the anonymous data segment or stack segment pages of a task.

By default, both kinds of memory spreading are off, and memory pages are allocated on the node local to where the task is running, except perhaps as modified by the tasks NUMA memory policy or cpuset configuration. This is true as long as sufficient free memory pages are available.

When new cpusets are created, they inherit the memory spread settings of their parent.

Setting memory spreading causes allocations for the affected page or slab caches to ignore the task's NUMA memory policy and be spread instead. Tasks using `mbind()` or `set_mempolicy()` calls to set NUMA memory policies will not notice any change in these calls, as a result of their containing tasks memory spread settings. If memory spreading is turned off, the currently specified NUMA memory policy once again applies to memory page allocations.

Both `cpuset.memory_spread_page` and `cpuset.memory_spread_slab` are Boolean flag files. By default, they contain 0. This means the feature is off for the cpuset. If a 1 is written to this file, the named feature is turned on for the cpuset.

This memory placement policy is also known (in other contexts) as round-robin or interleave.

This policy can provide substantial improvements for jobs that need to place thread local data on the corresponding node, but that need to access large file system data sets that need to be spread across the several nodes in the job's cpuset in order to fit. Without this policy, especially for jobs that might have one thread reading in the data set, the memory allocation across the nodes in the jobs cpuset can become very uneven.

## Memory Migration

Normally, under the default setting of `cpuset.memory_migrate`, once a page is allocated (given a physical page of main memory), that page stays on whatever node it was allocated, as long as it remains allocated, even if the cpuset's memory placement `mems` policy subsequently changes. The default setting has the `cpuset.memory_migrate` flag disabled.

When memory migration is enabled in a cpuset, if the `mems` setting of the cpuset is changed, any memory page in use by any task in the cpuset that is on a memory node no longer allowed is migrated to a memory node that is allowed.

Also, if a task is moved into a cpuset with `cpuset.memory_migrate` enabled, any memory pages it uses that were on memory nodes allowed in its previous cpuset, but which are not allowed in its new cpuset, are migrated to a memory node allowed in the new cpuset.

The relative placement of a migrated page within the cpuset is preserved during these migration operations if possible. For example, if the page was on the second valid node of the prior cpuset then the page will be placed on the second valid node of the new cpuset, if possible.

In order to maintain the cpuset relative position of pages, even pages on memory nodes allowed in both the old and new cpusets may be migrated. For example, if `cpuset.memory_migrate` is enabled in a cpuset, and that cpuset's `mems` file is written, changing it from memory nodes 4-7 to memory nodes 5-8, the following page migrations are done, in order, for all pages in the address space of tasks in that cpuset:

1. Migrate pages on node 7 to node 8.
2. Migrate pages on node 6 to node 7.
3. Migrate pages on node 5 to node 6.
4. Migrate pages on node 4 to node 5.

In this example, pages on any memory node other than 4 through 7 are not migrated. The order in which nodes are handled in a migration is intentionally chosen so as to avoid migrating memory to a node until any migrations from that node have first been accomplished.

## Mask Format

The mask format is used to represent CPU and memory node bitmasks in the `/proc/pid/status` file. It is hexadecimal, using ASCII characters 0 through 9 and a through f. This format displays each 32-bit word in hex (zero filled), and for masks longer than one word, uses a comma separator between words. Words are displayed in big endian order (most significant first). And hexadecimal digits within a word are also in big-endian order. The number of 32-bit words displayed is the minimum number needed to display all bits of the bitmask, based on the size of the bitmask. An example of the mask format is as follows:

00000001	# just bit 0 set
80000000,00000000,00000000	# just bit 95 set
00000001,00000000,00000000	# just bit 64 set
000000ff,00000000	# bits 32-39 set
00000000,000E3862	# bits 1,5,6,11-13,17-19 set

A mask with bits 0, 1, 2, 4, 8, 16, 32, and 64 set displays as 00000001,00000001,00010117. The first 1 is for bit 64, the second for bit 32, the third for bit 16, the fourth for bit 8, the fifth for bit 4, and the 7 is for bits 2, 1 and 0.

## List Format

The list format is used to represent CPU and memory node bitmasks (sets of CPU and memory node numbers) in the `/dev/cpuset` file system. It is a comma-separated list of CPU or memory node numbers and ranges of numbers, in ASCII decimal. An example of list format is as follows:

```
0-4,9          # bits 0, 1, 2, 3, 4, and 9 set
0-3,7,12-15    # bits 0, 1, 2, 3, 7, 12, 13, 14, and 15 set
```

## Cpuset Permissions

The permissions of a cpuset are determined by the permissions of the special files and directories in the cpuset file system, normally mounted at `/dev/cpuset`.

For example, a task can put itself in some other cpuset (than its current one) if it can write the `tasks` file (see "Cpuset File System Directories" on page 14) for that cpuset (requires execute permission on the encompassing directories and write permission on that `tasks` file).

An additional constraint is applied to requests to place some other task in a cpuset. One task may not attach another task to a cpuset unless it has permission to send that task a signal.

A task may create a child cpuset if it can access and write the parent cpuset directory. It can modify the CPUs or memory nodes in a cpuset if it can access that cpuset's directory (execute permissions on the encompassing directories) and write the corresponding `cpuset.cpus` or `cpuset.mems` file (see "Cpuset File System Directories" on page 14).

It should be noted, however, that changes to the CPUs of a cpuset do not apply to any task in that cpuset until the task is reattached to that cpuset. If a task can write the `cpuset.cpus` file, it should also be able to write the `tasks` file and might be expected to have permission to reattach the tasks therein (equivalent to permission to send them a signal).

There is one minor difference between the manner in which cpuset path permissions are evaluated by `libcgroup` and the manner in which file system operation permissions are evaluated by direct system calls. System calls that operate on file pathnames, such as the `open(2)` system call, rely on direct kernel support for a task's current directory. Therefore, such calls can successfully operate on files in or below a task's current directory, even if the task lacks search permission on some ancestor

directory. Calls in `libcputset` that operate on `cpuset` pathnames, such as the `cpuset_query()` call, rely on `libcputset` internal conversion of all `cpuset` pathnames to full, root-based paths. They cannot successfully operate on a `cpuset` unless the task has search permission on all ancestor directories, starting with the usual `cpuset` mount point (`/dev/cpuset`).

## CPU Scheduling and Memory Allocation for Cpusets

This section describes CPU scheduling and memory allocation for `cpusets` and covers these topics:

- "Linux Kernel CPU and Memory Placement Settings" on page 26
- "Manipulating Cpusets" on page 27

### Linux Kernel CPU and Memory Placement Settings

The Linux kernel exposes to user space three important attributes of each task that the kernel uses to control that tasks processor and memory placement, as follows:

- The `cpuset` path of each task, relative to the root of the `cpuset` file system, is available in the file `/proc/pid/cpuset`. For each task (PID), the file lists its `cpuset` path relative to the root of the `cpuset` file system.
- The actual CPU bitmask used by the kernel scheduler to determine on which CPUs a task may be scheduled is displayed in the `Cpus_allowed` field of the file `/proc/pid/status` for that task *pid*.
- The actual memory node bitmask used by the kernel memory allocator to determine on which memory nodes a task may obtain memory is displayed in the `Mems_allowed` field of the file of the file `/proc/pid/status` for that task *pid*.

Each of the above files is read-only. You can ask the kernel to make changes to these settings by using the various `cpuset` interfaces and the `sched_setaffinity(2)`, `mbind(2)`, and `set_mempolicy(2)` system calls.

The `cpus_allowed` and `mems_allowed` status file values for a task may differ from the `cpus` and `mems` values defined in the `cpuset` directory for the task for the following reasons:

- A task might call the `sched_setaffinity`, `mbind`, or `set_mempolicy` system calls to restrain its placement to less than its `cpuset`.

- Various temporary changes to `cpus_allowed` status file values are done by kernel internal code
- Attaching a task to a cpuset does not change its `mems_allowed` status file value until the next time that task needs kernel memory.
- Changing the CPUs in a cpuset does not change the `cpus_allowed` status file value of the tasks attached to the cpuset until those tasks are reattached to it (to avoid a hook in the hotpath scheduler code in the kernel).

Use the `cpuset_reattach` routine to perform this update after a changing the CPUs allowed to a cpuset.

- If hotplug is used to remove all the CPUs or all the memory nodes in a cpuset, the tasks attached to that cpuset will have their `cpus_allowed` status file values or `mems_allowed` status file values altered to the CPUs or memory nodes when the closest ancestor to that cpuset is not empty.

## Manipulating Cpusets

New cpusets are created using the `mkdir(1)` command (at the shell (see "Using Cpusets at the Shell Prompt" on page 27) or in C programs (see Appendix A, "Cpuset Library Functions" on page 45)). Old cpusets are removed using the `rmdir(1)` commands. The `Cpus_allowed` and `Mems_allowed` status file files are accessed using `read(2)` and `write(2)` system calls or shell commands such as `cat` and `echo`.

The CPUs and memory nodes in a given cpuset are always a subset of its parent. The root cpuset has all possible CPUs and memory nodes in the system. A cpuset may be exclusive (CPU or memory) only if its parent is similarly exclusive.

## Using Cpusets at the Shell Prompt

This section describes the use of cpusets using shell commands. For information on the `cpuset(1)` command line utility, see "Cpuset Command Line Utility" on page 29. For information on using the cpuset library functions, see Appendix A, "Cpuset Library Functions" on page 45.

When modifying the CPUs in a cpuset from the from the shell prompt, you must write the process ID (PID) of each task attached to that cpuset back into the cpuset's `tasks` file. When using the `libcpuset` API, use the `cpuset_reattach()` routine to

perform this step. The reasons for performing this step are described in "Modifying the CPUs in a Cpuset and Kernel Processing" on page 38.

The following example explains how to start a new job within a cpuset. You create a new cpuset called `green`, assign CPUs 2 and 3 and memory node 1 to the new cpuset, and start a subshell running in the cpuset. The steps are as follows:

1. The cpuset system is created and initialized by the kernel at system boot. You allow user space access to the cpuset system by mounting the cpuset virtual file system (VFS) at `/dev/cpuset`, as follows:

```
% mkdir /dev/cpuset  
% mount -t cpuset cpuset /dev/cpuset
```

---

**Note:** If the `mkdir(1)` and/or the `mount(8)` command fail, it is because they have already been run.

---

2. Type the following commands to create the new cpuset called `green` within the `/dev/cpuset` virtual file system:

```
% cd /dev/cpuset  
% mkdir green  
% cd green
```

3. Use the `echo` command to assign CPUs 2 and 3 and memory node 1 to the `green` cpuset, as follows:

```
% /bin/echo 2-3 > cpus  
% /bin/echo 1 > mems
```

4. Start a task to be the parent process of the new job, and attach the task to the new cpuset by writing its PID to the `/dev/cpuset/tasks` file for that cpuset.

For example:

```
% /bin/echo $$ > tasks  
sh
```

The subshell `sh` is now running in the `green` cpuset.



The file `/proc/self/cpuset` shows your current cpuset, as follows:

```
% cat /proc/self/cpuset
/green
```

From this shell, you can `fork`, `exec` or `clone(2)` the job tasks. By default, any child task of this shell will also be in cpuset `green`. To list the PIDs of the tasks currently in cpuset `green`, type the following:

```
% cat /dev/cpuset/green/tasks
4965
5043
```

In this example, PID 4965 is your shell, and PID 5043 is the `cat` command itself.

A cpuset must be empty before you remove it. To determine if you can remove the cpuset, take one of the following approaches:

- Run the `cat(1)` command on the cpuset directory `tasks` files to ensure no PIDs are listed. Or, within an application, use `libcpuset C API`.
- Run an `ls(1)` command on the cpuset directory to ensure it has no subdirectories.

For example if you want to remove the `green` cpuset, make sure the cpuset is empty. If the cpuset is not empty, when you remove it, the system generates a message similar to the following:

```
%rmdir green
rmdir: 'green': Device or resource busy
```

The following example explains how to remove cpuset `green` from the `/dev/cpuset` directory:

```
%cd /dev/cpuset
%rmdir green
```

## Cpuset Command Line Utility

The `cpuset(1)` command is used to create and destroy cpusets, to retrieve information about existing cpusets, and to attach processes to cpusets. The `cpuset(1)` command line utility is not essential to the use of cpusets. This utility provides an alternative that might be convenient for some uses.

A cpuset is defined by a cpuset configuration file and a name. For a definition of the cpuset configuration file format, see "Cpuset Text Format" on page 37.

The cpuset configuration file lists the CPUs and memory nodes that are members of the cpuset. It also contains any additional parameters required to define the cpuset. For more information on the cpuset configuration file, see "bootcpuset.conf File" on page 35.

This command automatically handles reattaching tasks to their cpuset whenever necessary, as described in the `cpuset_reattach` routine in Appendix A, "Cpuset Library Functions" on page 45.

The following are additional notes:

- The *csname* of / (slash) refers to the top cpuset, which encompasses all CPUs and memory nodes in the system. The *csname* of . (dot) refers to the cpuset of the current task. If a *csname* begins with the / (slash) character, it is resolved relative to the top cpuset, otherwise it is resolved relative to the cpuset of the current task.
- The *command input stream* and *command output stream* refer to the `stdin` (file descriptor 0) and `stdout` (file descriptor 1) of the command, unless the `-f` option is specified, in which case they refer to the file specified to `-f` option. Specifying the file name - to the `-f` option, as in `-f -`, is equivalent to not specifying the `-f` option at all.
- Specify exactly **one** of the action options. They are as follows:  
`-c, -m, -x, -d, -p, -a, -i, -w, -s, -R`
- You can specify additional modifier options in any order. All modifier options are evaluated first, before the action option. The modifier options are as follows:  
`-r, -I, -f`
- If you specify a help option, the command displays help information. No actions are run.

The following topics contain more information about the `cpusets(1)` command:

- "Action Options" on page 31
- "Modifier Options" on page 33
- "Help Options" on page 34

## Action Options

Specify exactly one of the following `cpuset(1)` command action options:

`-c csname, --create=csname`

Creates cpuset named *csname* using the cpuset text format (see "Cpuset Text Format" on page 37) representation read from the commands input stream.

`-m csname, --modify=csname`

Modifies the existing cpuset *csname* to have the properties in the cpuset text format (see "Cpuset Text Format" on page 37) representation read from the commands input stream.

`-x csname, --remove=csname`

Removes the cpuset named *csname*. A cpuset may only be removed if there are no processes currently attached to it and the cpuset has no descendant cpusets.

`-d csname, --dump=csname`

Writes a cpuset text format representation (see "Cpuset Text Format" on page 37) of the cpuset named *csname* to the commands output stream.

`-p csname, --procs=csname`

Lists to the commands output stream the processes (by `pid`) attached to the cpuset named *csname*. If the `-r` option is also specified, lists the `pid` of each process attached to any descendant of cpuset *csname*.

`-a csname, --attach=csname`

Attaches to the cpuset named *csname* the processes whose `pids` are read from the commands input stream, one `pid` per line.

`-i csname, --invoke=csname`

Invokes a command in the cpuset named *csname*. If `-I` option is set, use that command and arguments, otherwise if the environment variable `$SHELL` is set, use that command, otherwise, use `/bin/sh`.

`-w pid, --which=pid`

Lists the name of the cpuset to which process *pid* is attached, to the commands output stream. If *pid* is zero (0), then the full cpuset path of the current task is displayed.

`-s csname, --show=csname`

Prints to the commands output stream the names of the cpusets below cpuset *csname*. If the `-r` option is also specified, this recursively includes *csname* and all its descendants, otherwise it just includes the immediate child cpusets of *csname*. The cpuset names are printed one per line.

`-R csname, --reattach=r`

Reattaches each task in cpuset *csname*. This is required after changing the `cpus` value of a cpuset, in order to get the tasks already attached to that cpuset to rebind to the changed CPU placement.

`-z csname, --size=csname`

Prints the size of (number of CPUs in) a cpuset to the commands output stream, as an ASCII decimal newline terminated string.

`-F flist, --family=flist`

Creates a family of non-overlapping child cpusets, given an *flist* of cpuset names and sizes (number of CPUs). Fails if the total sizes exceeds the size of the current cpuset. Enter cpuset names relative to the current cpuset, and their requested size, as alternating command line arguments. For example:

```
cpuset -F foo 2 bar 6 baz 4
```

This creates three child cpusets named `foo`, `bar`, and `baz`, having 2, 6, and 4 CPUs, respectively.

This example will fail with an error message and a nonzero exit status if the current cpuset lacks at least 12 CPUs.

These cpuset names are relative to the current cpuset and will not collide with the cpuset names descendent from other cpusets. Hence

two commands, running in different cpusets, can both create a child cpuset named `foo` without a problem.

## Modifier Options

You can specify modifier options in any combination. The modifier options are as follows:

`-r, --recursive`

When used with `-p` or `-s` option, applies to all descendants recursively of the named cpuset *csname*.

`-I cmd, --invokecmd=cmd`

When used with the `-i` option, the command *cmd* is invoked, with any optional unused arguments. The following example invokes an interactive subshell in cpuset `foo`:

```
cpuset -i foo -I sh -- -i
```

The next example invokes a second cpuset command in cpuset `foo`, which then displays the full cpuset path of `foo`:

```
cpuset -i foo -I cpuset -- -w 0
```

---

**Note:** The double minus `--` is needed to end option parsing by the initial cpuset command.

---

`-f fname, --file=fname`

Uses file named *fname* for command input or output stream, instead of `stdin` or `stdout`.

`--move_tasks_from=csname1 --move_tasks_to=csname2`

Move all tasks from cpuset *csname1* to cpuset *csname2*. Retries up to ten times to move all tasks, in case it is racing against parallel attempts to fork or add tasks into cpuset *csname1*. Fails with nonzero

exit status and an error message to `stderr` if unable to move all tasks out of *csname1*.

## Help Options

The help options override all other options and are as follows:

`-h, --help`

Displays command usage

## Boot Cpuset

You can use the `bootcpuset(8)` command to create a boot cpuset during the system boot that you can use to restrict the default placement of almost all UNIX processes on your system. You can use the `bootcpuset` to reduce the interference of system processes with applications running on dedicated cpusets.

The default cpuset for the `init` process, classic UNIX daemons, and user login shells is the root cpuset that contains the entire system. For systems dedicated to running particular applications, it is better to restrict `init`, the kernel daemons, and login shells to a particular set of CPUs and memory nodes called the `bootcpuset`.

This section covers the following topics:

- "Creating a Bootcpuset" on page 34
- "`bootcpuset.conf` File" on page 35

## Creating a Bootcpuset

This section describes how to create a `bootcpuset`.

**Procedure 2-3** To create a `bootcpuset`

1. Create a `/etc/bootcpuset.conf` file with values to restrict system processes to the CPUs and memory nodes appropriate for your system, similar to the following:

```
cpus 0-7
mems 0
```

2. Update the boot file to append `init=/sbin/bootcpuset` to the kernel command line.

Depending on the operating system, you might need to update `grub.conf`, `grub.cfg`, or `elilo.conf`. For more information, see the operating system documentation.

3. Reboot the system.

Subsequent system reboots restrict most processes to the bootcpuset defined in the boot files.

### **bootcpuset.conf File**

The `/etc/bootcpuset.conf` file describes what CPUs and memory nodes are to be in the bootcpuset. The kernel boot command line option `init` is used to invoke the `/sbin/bootcpuset` binary ahead of the `/sbin/init` binary, using the `elilo` syntax: `append="init=/sbin/bootcpuset"`.

When invoked with `pid=1`, the `/sbin/bootcpuset` binary does the following:

- Sets up a bootcpuset (configuration defined in the `/etc/bootcpuset.conf` file).
- Attaches itself to this bootcpuset.
- Attaches any unpinned kernel threads to it.
- Invokes an `exec` call to execute `/sbin/init`, `/etc/init`, `/bin/init` or `/bin/sh`.

A kernel thread is deemed to be unpinned (third bullet in the list above) if its `Cpus_allowed` value (as listed in that thread's `/proc/pid/status` file for the `Cpus_allowed` field) allows running on all online CPUs. Kernel threads that are restricted to some proper subset of CPUs are left untouched, under the assumption that they have a good reason to be running on those restricted CPUs. Such kernel threads as migration (to handle moving threads between CPUs) and `ksoftirqd` (to handle per-CPU work off interrupts) must be pinned to each CPU or each memory node.

Comments in the `/etc/bootcpuset.conf` configuration file begin with the pound (`#`) character and extend to the end of the line. After stripping comments, the bootcpuset command examines the first white space separated token on each line.

If the first token on the line matches `mems` or `mem` (case insensitive match) then the second token on the line is written to the `/dev/cpuset/boot/cpuset.mems` file.

If the first token on the line matches `cpus` or `cpu` (case insensitive match), then the second token is written to the `/dev/cpuset/boot/cpuset.cpus` file.

If the first token in its entirety matches (case insensitive match) `"verbose"`, the `bootcpuset` command prints a trace of its actions to the console. A typical such trace has 20 or 30 lines, detailing the steps taken by `/sbin/bootcpuset` and is useful in understanding its behavior and analyzing problems. The `bootcpuset` command ignores all other lines in the `/etc/bootcpuset.conf` configuration file.

## Configuring a User Cpuset for Interactive Sessions

You might want to create the following two cpusets for interactive sessions:

1. A *boot* cpuset, which runs the traditional daemon and server tasks. Create this cpuset first.
2. A *user* cpuset, which can hold interactive `telnet`, `rlogin` and/or secure shell (SSH) user sessions.

When you create a user cpuset, you provide additional isolation between interactive user login sessions and essential system tasks. For example, if a user process in the user cpuset consumes excessive CPU, system file buffer cache, or memory resources, the user cpuset contains the process so it does not seriously affect essential system services in the boot cpuset.

HPE provides example scripts that you can use to create a user cpuset called `user`. The `user` cpuset runs as one of the last scripts when the system boots. The sample script requires that you create a boot cpuset. For instructions that explain how to create the boot cpuset, see the script that is appropriate for your operating system, as follows:

- For RHEL 7 and SLES 12 operating systems, use the script in the following location:

```
/usr/share/cpuset/samples/sysd_usercpuset
```

- For RHEL 6 and SLES 11 operating systems, use the script in the following location:

```
/usr/share/cpuset/samples/usercpuset
```



The script puts the `sshd` and `xinetd` daemons that create interactive login sessions into this new user cpuset. The user cpuset configuration is defined in the `/etc/usercpuset.conf` file.

To isolate the boot cpuset from the user cpuset, make sure that the set of `cpus` and `mems` values in the `/etc/bootcpuset.conf` file does not overlap the `cpus` and `mems` values in the `/etc/usercpuset.conf` file. Instructions for using each script are included in the comments within the script.

## Cpuset Text Format

Cpuset settings may be exported to and imported from text files using a text format representation of cpusets.

Permissions of files holding these text representations have no special significance to the implementation of cpusets. Rather, the permissions of the special cpuset files in the cpuset file system, normally mounted at `/dev/cpuset`, control reading and writing of and attaching to cpusets.

The text representation of cpusets is not essential to the use of cpusets. One can directly manipulate the special files in the cpuset file system. This text representation provides an alternative that may be convenient for some uses and a form for representing cpusets that users of earlier versions of cpusets will find familiar.

The exported cpuset text format has fewer directives than earlier Linux versions. Additional directives may be added in the future.

The cpuset text format supports one directive per line. Comments begin with the pound character (`#`) and extend to the end of line.

After stripping comments, the first white space separated token on each remaining line selects from the following possible directives:

<code>cpus</code>	Specifies which CPUs are in this cpuset. The second token on the line must be a comma-separated list of CPU numbers and ranges of numbers.
<code>mems</code>	Specify which memory nodes are in this cpuset. The second token on the line must be a comma-separated list of memory node numbers and ranges of numbers.
<code>cpu_exclusive</code>	The <code>cpu_exclusive</code> flag is set.
<code>mem_exclusive</code>	The <code>mem_exclusive</code> flag is set.

`notify_on_release`      The `notify_on_release` flag is set

Additional unnecessary tokens on a line are quietly ignored. Lines containing only comments and white space are ignored.

The token `cpu` is allowed for `cpus` and `mem` for `mems`. Matching is case insensitive.

See the `libcputset` routines `cpuset_import` and `cpuset_export` to handle converting the internal `struct cpuset` representation of cpusets to (export) and from (import) this text representation.

For information on manipulating `cpuset` text files at the shell prompt or in shell scripts using the `cpuset(1)` command, see "Cpuset Command Line Utility" on page 29.

## Modifying the CPUs in a Cpuset and Kernel Processing

In order to minimize the impact of cpusets on critical kernel code, such as the scheduler, and due to the fact that the Linux kernel does not support one task updating the memory placement of another task directly, the impact on a task of changing its `cpuset` CPU or memory node placement or of changing to which `cpuset` a task is attached, is subtle and is described in the following paragraphs.

When a `cpuset` has its memory nodes modified, for each task attached to that `cpuset`, the next time that the kernel attempts to allocate a page of memory for a particular task, the kernel notices the change in the task's `cpuset`, and updates its per-task memory placement to remain within the new `cpusets` memory placement. If the task was using memory policy `MPOL_BIND` and the nodes to which it was bound overlaps with its new `cpuset`, the task continues to use whatever subset of `MPOL_BIND` nodes that are still allowed in the new `cpuset`. If the task was using `MPOL_BIND` and now none of its `MPOL_BIND` nodes are allowed in the new `cpuset`, the task is essentially treated as if it was `MPOL_BIND` bound to the new `cpuset` (even though its NUMA placement, as queried by the `get_mempolicy()` routine, does not change). If a task is moved from one `cpuset` to another, the kernel adjusts the task's memory placement, as above, the next time that the kernel attempts to allocate a page of memory for that task.

When a `cpuset` has its CPUs modified, each task using that `cpuset` **does not change** its behavior automatically. In order to minimize the impact on the critical kernel scheduling code, tasks continue to use their prior CPU placement until they are rebound to their `cpuset` by rewriting their PID to the `tasks` file of their `cpuset`. If a task is moved from one `cpuset` to another, its CPU placement is updated in the same way as if the task's PID is rewritten to the `tasks` file of its current `cpuset`.

In summary, the memory placement of a task whose cpuset is changed is automatically updated by the kernel, on the next allocation of a page for that task but the processor placement is not updated until that task's PID is rewritten to the `tasks` file of its cpuset. The delay in rebinding a task's memory placement is necessary because the kernel does not support one task changing memory placement of another task. The added user level step in rebinding a task's CPU placement is necessary to avoid impacting the scheduler code in the kernel with a check for changes in a task's processor placement.

## Using Cpusets with Hyper-Threads

*Threading* in a software application splits instructions into multiple streams so that multiple processors can act on them.

Hyper-Threading (HT) Technology, developed by Intel Corporation, provides thread-level parallelism on each processor, resulting in more efficient use of processor resources, higher processing throughput, and improved performance. One physical CPU can appear as two logical CPUs by having additional registers to overlap two instruction streams or a single processor can have dual cores executing instructions in parallel.

In addition to their traditional use to control the placement of jobs on the CPUs and memory nodes of a system, cpusets also provide a convenient mechanism to control the use of Hyper-Threading Technology.

Some jobs achieve better performance by using both of the Hyper-Threaded sides, A and B, of a processor core, and some run better by using just one of the sides, allowing the other side to idle.

Since each logical (Hyper-Threaded) processor in a core has a distinct CPU number, you can specify a cpuset that contains both sides of a processor core or a cpuset that contains just one side from a processor core.

Cpusets can be configured to include any combination of the logical CPUs in a system.

For example, the following cpuset configuration file called `cpuset.cfg` includes the A sides of an HT enabled system, along with all the memory, on the first 32 nodes (assuming 2 cores per node). The colon (:) prefixes the stride. The stride of 2 in this example means use every other logical CPU.

```
cpus 0-127:2    # the even numbered CPUs 0, 2, 4, ... 126
mems 0-63      # all memory nodes 0, 1, 2, ... 63
```

To create a cpuset called `foo` and run a job called `bar` in that cpuset, defined by the cpuset configuration file `cpuset.cfg` shown above, use the following commands:

```
cpuset -c /foo < cpuset.cfg
cpuset -i /foo -I bar
```

To specify both sides of the first 64 cores, use the following entry in your cpuset configuration file:

```
cpus 0-127
```

To specify just the B sides of the processor cores of an HT enabled system, use the following entry in your cpuset configuration file:

```
cpus 1-127:2
```

The examples above assume that the CPUs are uniformly numbered with the even numbers for the A side and odd numbers for the B side of the processors cores. This is usually the case, but not guaranteed. You can still place a job on a system that is not uniformly numbered. Currently, it involves a longer argument list to the `cpus` option, that is, you must explicitly list the desired CPUs.

If you are using a `bootcpuset` to keep other tasks confined, you do not need to create a separate cpuset with just the B side CPUs to avoid having some tasks running on the B sides of the processor cores. If there is no cpuset for the B sides of the processor cores, except the all encompassing root cpuset, and if only root can put tasks in the root cpuset, then no one other tasks can run on the B sides.

You can use the `dplace(1)` command to manage more detailed placement of job tasks within a cpuset. Since the `dplace` command numbering of CPUs is relative to the cpuset, it does not affect the `dplace` configuration. This is true in the case where the cpuset includes both sides of Hyper-Threaded cores, just one side of the Hyper-Threaded cores, or even is on a system that does not support Hyper-threading.

Typically, the logical numbering of CPUs puts the even numbered CPUs on the A sides of processor cores and the odd numbered CPUs on the B sides. You can easily specify that only every other side is used using the stride suffix `:2`, described above. If the CPU number range starts with an even number, the A sides of the processor cores are used. If the CPU range starts with an odd number, the B sides of the processor cores are used.

The following procedure sets up a job to run only on the A sides of the system's Hyper-Threaded cores and to ensure that no other tasks run on the B sides, which remain idle.

**Procedure 2-4** To Configure a System with Hyper-Threaded Cores

1. Verify that you have root (sudo or superuser) permission.
2. Define a bootcpuset to restrain the kernel, system daemon, and user login session threads to a designated set of CPUs.
3. Define a cpuset that includes the A sides of the processors to be used for this job.

Make sure no cpuset is created using the B side CPUs in these processors. This precaution prevents disruptive tasks from running on the corresponding B side CPUs.

If you use a bootcpuset to confine the traditional UNIX load processes, nothing will run on the other CPUs in the system, except when those CPUs are included in a cpuset to which a job has been assigned. These CPUs are of course in the root cpuset, however, this cpuset is normally only usable by a system administrator or batch scheduler with root permissions. This prevents any user without root permission from running a task on those CPUs, unless an administrator or service with root permission allows it. For more information, see "Boot Cpuset" on page 34.

A `ps(1)` or `top(1)` invocation will show a handful of threads on unused CPUs. These are kernel threads assigned to every CPU in support of user applications running on those CPUs to handle tasks such as asynchronous file system writes and task migration between CPUs. If no application is actually using a CPU, the kernel threads on that CPU will be almost always idle.

## Cpuset Programming Model

Cpusets use a hierarchical name space. The name space remains visible to all tasks on a system. Once created, a cpuset remains in existence until it is deleted or until the system is rebooted, even if no tasks are currently running in that cpuset.

The key properties of a cpuset are its pathname, the list of which CPUs and memory nodes it contains, and whether the cpuset has exclusive rights to these resources.

Every task (process) in the system is attached to (running inside) a cpuset. Tasks inherit their parents cpuset attachment when forked. This binding of task to a cpuset can subsequently be changed, either by the task itself, or externally from another task, given sufficient authority.

Tasks have their CPU and memory placement constrained to whatever their containing cpuset allows. A cpuset may have exclusive rights to its CPUs and memory, which provides certain guarantees that other cpusets will not overlap.

At system boot, a top level root cpuset is created, which includes all CPUs and memory nodes on the system. The usual mount point of the cpuset file system and therefore the usual file system path to this root cpuset, is `/dev/cpuset`.

Optionally, a "boot" cpuset may be created, at `/dev/cpuset/boot`, to include typically just a one or a few CPUs and memory nodes. A typical use for a "boot" cpuset is to contain the general purpose UNIX daemons and login sessions, while reserving the rest of the system for running specific major applications on dedicated cpusets. For more information, see "Boot Cpuset" on page 34.

Moved tasks do not have the memory they might have allocated on their old nodes moved to the new nodes. On kernels that support such memory migration, use the [optional] `cpuset_migrate` to move allocated memory as well.

Cpusets have a permission structure which determines which users have rights to query, modify, and attach to any given cpuset. Rights are based on the hierarchical model provided by the underlying Linux 2.6 kernel cpuset file system.

To create a cpuset from within a C language application, your program obtains a handle to a new `struct cpuset`, sets the desired attributes via that handle, and issues a `cpuset_create()` call to create the desired cpuset and bind it to the specified name. Your program can also issue calls to list by name what cpusets exist, query their properties, move tasks between cpusets, list what tasks are currently attached to a cpuset, and delete cpusets.

The names of cpusets in this C library are always relative to the root cpuset mount point, typically `/dev/cpuset`. For more information on the `libcputset` C language application programming interface (API) functions, see Appendix A, "Cpuset Library Functions" on page 45.

## System Error Messages

The Linux kernel implementation of cpusets sets `errno` to specify the reason for a failed system call that affects cpusets. These `errno` values are available when a cpuset library call fails. They can be displayed by shell commands used to directly manipulate files below the `/dev/cpuset` directory and can be displayed by the `cpuset(1)` command.

The possible `errno` settings and their meaning when set on a failed cpuset call are, as follows:

ENOSYS	Invoked on an operating system kernel that does not support cpusets.
ENODEV	Invoked on a system that supports cpusets, but when the cpuset file system is not currently mounted at <code>/dev/cpuset</code> .
ENOMEM	Insufficient memory is available.
EBUSY	Attempted <code>cpuset_delete()</code> on a cpuset with attached tasks.
EBUSY	Attempted <code>cpuset_delete()</code> on a cpuset with child cpusets.
ENOENT	Attempted <code>cpuset_create()</code> in a parent cpuset that does not exist.
EEXIST	Attempted <code>cpuset_create()</code> for a cpuset that already exists.
E2BIG	Attempted a <code>write(2)</code> system call on a special cpuset file with a length larger than some kernel determined upper limit on the length of such writes.
ESRCH	Attempted to <code>cpuset_move()</code> a nonexistent task.
EACCES	Attempted to <code>cpuset_move()</code> a task that the process lacks permission to move.
ENOSPC	Attempted to <code>cpuset_move()</code> a task to an empty cpuset.
EINVAL	The <code>relcpu</code> argument to <code>cpuset_pin()</code> function is out of range (not between "zero" and " <code>cpuset_size() - 1</code> ").
EINVAL	Attempted to change a cpuset in a way that would violate a <code>cpu_exclusive</code> or <code>mem_exclusive</code> attribute of that cpuset or any of its siblings.
EINVAL	Attempted to write an empty <code>cpus</code> or <code>mems</code> bitmask to the kernel. The kernel creates new cpusets (using the <code>mkdir</code> function) with empty <code>cpuset.cpus</code> and <code>cpuset.mems</code> files and the user level cpuset and bitmask code works with empty masks. But the kernel

	will not allow an empty bitmask (no bits set) to be written to the special <code>cpus</code> or <code>mems</code> files of a <code>cpuset</code> .
EIO	Attempted to <code>write(2)</code> a string to a <code>cpuset</code> tasks file that does not begin with an ASCII decimal integer.
ENOSPC	Attempted to <code>write(2)</code> a list to a <code>cpuset.cpus</code> file that did not include any online CPUs.
ENOSPC	Attempted to <code>write(2)</code> a list to a <code>cpuset.mems</code> file that did not include any online memory nodes.
EACCES	Attempted to add a CPU or memory resource to a <code>cpuset</code> that is not already in its parent.
EACCES	Attempted to set the <code>cpu_exclusive</code> or <code>mem_exclusive</code> flag on a <code>cpuset</code> whose parent lacks the same setting.
EBUSY	Attempted to remove a CPU or memory resource from a <code>cpuset</code> that is also in a child of that <code>cpuset</code> .
EFAULT	Attempted to read or write a <code>cpuset</code> file using a buffer that was outside your accessible address space.
ENAMETOOLONG	Attempted to read a <code>/proc/<i>pid</i>/cpuset</code> file for a <code>cpuset</code> path that was longer than the kernel page size.



## Cpuset Library Functions

This appendix describes the `libcpuset` C application programming interface (API) functions and covers the following topics:

- "Basic Cpuset Library Functions" on page 45
- "Extensible Application Programming Interface" on page 46
- "Advanced Cpuset Library Functions" on page 47

### Basic Cpuset Library Functions

The basic `cpuset` C API supports the following functions:

- `cpuset_pin`, which pins the current thread to a CPU, preferring local memory.
- `cpuset_size`, which returns the number of CPUs that are in the current tasks `cpuset`.
- `cpuset_unpin`, which removes the effect of `cpuset_pin` and lets the task have run of its entire `cpuset`.
- `cpuset_where`, which returns on which CPU in current tasks `cpuset` did the task most recently execute.
- `cpuset_function`, which returns either a pointer to the named `libcpuset.so` function or `NULL`. For [optional] `cpuset` functions, `cpuset_function` enables dynamic adaptation to runtime environments that might or might not support a particular function.
- `cpuset_version`, which returns the `libcpuset` version number. This is an [optional] `cpuset` function.

For more information about the basic `cpuset` functions, see the `libcpuset(3)` man page.

## Extensible Application Programming Interface

In order to provide for the convenient and robust extensibility of this cpuset API over time, the following function enables dynamically obtaining pointers for optional functions by name, at runtime:

```
void *cpuset_function(const char * function_name)
```

It returns a function pointer or NULL if *function\_name* is not recognized.

For maximum portability, you should not reference any optional cpuset function by explicit name.

However, if you presume that an optional function will always be available on the target systems of interest, you might decide to explicitly reference it by name, in order to improve the clarity and simplicity of the software in question.

Also to support robust extensibility, flags and integer option values have names dynamically resolved at runtime, not via preprocessor macros.

Some functions in Advanced Cpuset Library Functions are marked [optional]. (see "Advanced Cpuset Library Functions" on page 47). They are not available in all implementations of libcpuset. Additional [optional] cpuset\_\* functions may also be added in the future. Functions that are not marked [optional] are available on all implementations of libcpuset.so and can be called directly without using cpuset\_function(). However, any of them can also be called indirectly via cpuset\_function().

To safely invoke an optional function, such as for example cpuset\_migrate(), use the following call sequence:

```
/*                                                                    */
/* NOTE: The example shown is for illustrative purposes only and */
/* has not been evaluated for use in a production environment. */
/*                                                                    */
/* fp has function signature of pointer to cpuset_migrate() */
int (*fp)(struct cpuset *fromcp, struct cpuset *to_cp, pid_t pid);
fp = cpuset_function("cpuset_migrate");
if (fp) {
    fp( ... );
} else {
    puts ("cpuset migration not supported");
}
```



---

**Caution:** The preceding example is for illustrative purposes only and has not been evaluated for use in a production environment.

---

If you invoke an [optional] function directly, your resulting program will not be able to link with any version of `libcpuset.so` that does not define that particular function.

## Advanced Cpuset Library Functions

The advanced cpuset API provides functions usable from a C language application for managing cpusets on a system-wide basis.

These functions primarily deal with the following three entities:

- `struct cpuset * structure`
- System cpusets
- Tasks

The `struct cpuset * structure` provides a transient in-memory structure used to build up a description of an existing or desired cpuset. These structures can be allocated, freed, queried, and modified.

Actual kernel cpusets are created under the `/dev/cpuset` directory, which is the usual mount point of the kernel's virtual cpuset filesystem. These cpusets are visible to all tasks in the system, and persist until the system is rebooted or until the cpuset is explicitly deleted. These cpusets can be created, deleted, queried, modified, listed, and examined.

Every task (also known as a process) is bound to exactly one cpuset at a time. You can list which tasks are bound to a given cpuset, and to which cpuset a given task is bound. You can change to which cpuset a task is bound.

The primary attributes of a cpuset are its lists of CPUs and memory nodes. The scheduling affinity for each task, whether set by default or explicitly by the `sched_setaffinity()` system call, is constrained to those CPUs that are available in that task's cpuset. The NUMA memory placement for each task, whether set by default or explicitly by the `mbind()` system call, is constrained to those memory nodes that are available in that task's cpuset. This provides the essential purpose of cpusets - to constrain the CPU and Memory usage of tasks to specified subsets of the system.

The other essential attribute of a cpuset is its pathname beneath `/dev/cpuset`. All tasks bound to the same cpuset pathname can be managed as a unit, and this hierarchical name space describes the nested resource management and hierarchical permission space supported by cpusets. Also, this hierarchy is used to enforce strict exclusion, using the following rules:

- A cpuset may only be marked strictly exclusive for CPU or memory if its parent is also.
- A cpuset may not make any CPUs or memory nodes available that are not also available in its parent.
- If a cpuset is exclusive for CPU or memory, it may not overlap CPUs or memory with any of its siblings.

The combination of these rules enables checking for strict exclusion just by making various checks on the parent, siblings, and existing child cpusets of the cpuset being changed, without having to check all cpusets in the system.

On error, some of these routines return -1 or NULL and set `errno`. If one of the routines below that requires cpuset kernel support or the cpuset file system mounted is invoked on an operating system kernel that does not support cpusets, then that routine returns failure and `errno` is set to `ENOSYS`. If invoked on a system that supports cpusets, but when the cpuset file system is not currently mounted at `/dev/cpuset`, it returns failure and `errno` is set to `ENODEV`.

The following inclusion and linkage provides access to the cpuset API from C code:

```
/*                                                    */
/* NOTE: The example shown is for illustrative purposes only and */
/* has not been evaluated for use in a production environment. */
/*                                                    */
#include <bitmask.h>
#include <cpuset.h>
/* link with -lcpuset */
```



**Caution:** The preceding example is for illustrative purposes only and has not been evaluated for use in a production environment.

---

**Note:** Some functions are marked [optional]. For an explanation, see "Extensible Application Programming Interface" on page 46.

---

The following functions are supported in the advanced cpuset C API:

**Allocate and free struct cpuset \* structure**

- `cpuset_alloc` - Returns handle to newly allocated struct cpuset \* structure
- `cpuset_free` - Discards no longer needed struct cpuset \* structure

**Lengths of CPUs and memory nodes bitmasks - needed to allocate bitmasks**

- `cpuset_cpus_nbits` - Number of bits needed for a CPU bitmask on current system
- `cpuset_mems_nbits` - Number of bits needed for a memory bitmask on current system

**Set various attributes of a struct cpuset \* Structure**

- `cpuset_setcpus` - Specifies CPUs in cpuset
- `cpuset_setmems` - Specifies memory nodes in cpuset
- `cpuset_set_iopt` - Specifies an integer value option of cpuset
- `cpuset_set_sopt` - [optional] Specifies a string value option of cpuset

**Query various attributes of a struct cpuset \* Structure**

- `cpuset_getcpus` - Queries CPUs in cpuset
- `cpuset_getmems` - Queries memory nodes in cpuset
- `cpuset_cpus_weight` - Number of CPUs in a cpuset
- `cpuset_mems_weight` - Number of memory nodes in a cpuset
- `cpuset_get_iopt` - Query an integer value option of cpuset
- `cpuset_get_sopt` - [optional] Species a string value option of cpuset

**Local CPUs and memory nodes**

- `cpuset_localcpus` - Queries the CPUs local to specified memory nodes
- `cpuset_localmems` - Queries the memory nodes local to specified CPUs
- `cpuset_cpumemdist` - [optional] Hardware distance from CPU to memory node
- `cpuset_cpu2node` - Returns number of memory node closed to specified CPU

- `cpuset_addr2node` - Return number of memory node holding page at specified address.
- `cpuset_cpu2node_c_rel` - Returns cpuset relative number of memory node closest to the specified CPU

**Create, delete, query, modify, list, and examine cpusets**

- `cpuset_create` - Creates a named cpuset as specified by `struct cpuset *` structure
- `cpuset_delete` - Deletes the specified cpuset (if empty)
- `cpuset_query` - Sets the `struct cpuset` structure to settings of specified cpuset
- `cpuset_modify` - Modifies the settings of a cpuset to those specified in a `struct cpuset` structure
- `cpuset_getcpusetpath` - Gets path of a tasks (0 for current) cpuset
- `cpuset_cpusetofpid` - Sets the `struct cpuset` structure to settings of cpuset of specified task
- `cpuset_mountpoint` - Returns path at which cpuset filesystem is mounted
- `cpuset_collides_exclusive` - [optional] True, if it would collide an exclusive
- `cpuset_nuke` - Removes cpuset any way possible

**List tasks (pids) currently attached to a cpuset**

- `cpuset_init_pidlist` - Initializes a list of tasks (pids) attached to a cpuset
- `cpuset_pidlist_length` - Returns number of elements in a list of pid
- `cpuset_get_pidlist` - Returns i'th element of a list of pids
- `cpuset_freepidlist` - Deallocates a list of pids

**Attach tasks to cpusets**

- `cpuset_move` - Moves task (0 for current) to a cpuset
- `cpuset_move_all` - Moves all tasks in a list of pids to a cpuset
- `cpuset_migrate` - [optional] Moves a task and its memory to a cpuset

- `cpuset_migrate_all` - [optional] Moves all tasks with memory in a list of pids to a cpuset
- `cpuset_reattach` - Rebinds `cpus_allowed` of each task in a cpuset after changing its cpus
- `cpuset_move_cpuset_tasks` - [optional] Moves all tasks in a cpuset to another cpuset

**Determine memory pressure**

- `cpuset_open_memory_pressure` - [optional] Opens handle to read memory\_pressure
- `cpuset_read_memory_pressure` - [optional] Reads cpuset current memory\_pressure
- `cpuset_close_memory_pressure` - [optional] Closes handle to read memory\_pressure

**Map between cpuset relative and system-wide CPU and memory node numbers**

- `cpuset_c_rel_to_sys_cpu` - Maps cpuset relative CPU number to system wide number
- `cpuset_c_sys_to_rel_cpu` - Maps system-wide CPU number to cpuset relative number
- `cpuset_c_rel_to_sys_mem` - Maps cpuset relative memory node number to system wide number
- `cpuset_c_sys_to_rel_mem` - Maps system-wide memory node number to cpuset relative number
- `cpuset_p_rel_to_sys_cpu` - Maps task cpuset relative CPU number to system wide number
- `cpuset_p_sys_to_rel_cpu` - Maps system-wide CPU number to task cpuset relative number
- `cpuset_p_rel_to_sys_mem` - Maps task cpuset relative memory node number to system-wide number
- `cpuset_p_sys_to_rel_mem` - Maps system-wide memory node number to task cpuset relative number

### **Placement operations for detecting cpuset migration**

- `cpuset_get_placement` - [optional] Returns the current placement of task `pid`
- `cpuset_equal_placement` - [optional] True, if two placements are equal
- `cpuset_free_placement` - [optional] Free placement

### **Bind to a CPU or memory node within the current cpuset**

- `cpuset_cpubind` - Binds to a single CPU within a cpuset (uses `sched_setaffinity(2)`)
- `cpuset_latestcpu` - Most recent CPU on which a task has executed
- `cpuset_membind` - Binds to a single memory node within a cpuset (uses `set_mempolicy(2)`)

### **Export cpuset settings to a regular file and import them from a regular file**

- `cpuset_export` - Exports cpuset settings to a text file
- `cpuset_import` - Imports cpuset settings from a text file

### **Traverse a cpuset hierarchy**

- `cpuset_fts_open` - Opens a cpuset hierarchy
- `cpuset_fts_read` - Obtains the next entry in the hierarchy
- `cpuset_fts_reverse` - Reverses the order of the cpusets in the hierarchy
- `cpuset_fts_rewind` - Rewinds to the first cpuset in the list
- `cpuset_fts_get_path` - Returns the cpuset path of an entry
- `cpuset_fts_get_stat` - Returns the `stat(2)` pointer information about an entry
- `cpuset_fts_get_cpuset` - Returns the cpuset pointer of an entry
- `cpuset_fts_get_errno` - Returns the `err` field of an entry
- `cpuset_fts_get_info` - Returns the `info` field of an entry
- `cpuset_fts_close` - Closes a cpuset hierarchy

### **Cpuset Library Functions Calling Sequence**



A typical calling sequence would use the above functions in the following order to create a new cpuset named `xyz` and attach itself to it, as follows:

```
/*                                                                    */
/* NOTE: The example shown is for illustrative purposes only and */
/* has not been evaluated for use in a production environment. */
/*                                                                    */
struct cpuset *cp = cpuset_alloc();
various cpuset_set*(cp, ...) calls
cpuset_create(cp, "xyz");
cpuset_free(cp);
cpuset_move(0, "xyz");
```



**Caution:** The preceding example is for illustrative purposes only and has not been evaluated for use in a production environment.

---



---

## Index

### C

#### Cpuset Facility

- advantages, 1
- boot cpuset, 34
  - creating, 34
  - /etc/bootcpuset.conf file, 35
- command line utility, 29
- cpuset
  - definition, 1
  - determine if cpusets are installed, 8
  - errno settings, 43
  - modifying CPUs and kernel processing, 38
  - system error messages, 42
- cpuset permissions, 25
- cpuset text format, 37
- directories, 14

- overview, 1
- programming model, 41
- scheduling and memory allocation, 26
- systems calls
  - mbind, 1
  - sched\_setaffinity, 1
  - set\_mempolicy, 1
- using cpusets at shell prompt, 27
  - create a cpuset, 27
  - remove a cpuset, 29
- Creating a cpuset, 27

### R

- Removing a cpuset, 29